
Multi-Agent Coordination via Support-Preserving Distillation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Offline MARL increasingly relies on generative policies to model multimodal joint
2 behavior, typically by distilling a centralized teacher into decentralized one-step
3 actors under the CTDE. We identify a failure mode at the teacher training stage:
4 standard flow-based teachers pair noise with replay targets independently, so nearby
5 noise samples can be routed toward conflicting coordination modes. The teacher
6 then produces samples between valid modes, and because the distillation loss
7 regresses each local actor onto the conditional mean of the teacher’s output given
8 local input, this error is not absorbed but propagated to the student. To remove this
9 teacher-side artifact, we propose *Mode-Support Semi-Discrete Optimal Transport*
10 (MoSDOT), which summarizes multimodal replay into a finite mode support with
11 prescribed capacities and uses conditional semi-discrete optimal transport to assign
12 each noise sample to a single mode before teacher training. We additionally study
13 a common-randomness variant that uses a shared noise component at execution to
14 expose the residual gap intrinsic to strict-product execution. On controlled diag-
15 nostics and offline MARL benchmarks, MoSDOT improves endpoint quality and
16 routing consistency, particularly on datasets exhibiting multimodal joint behavior.

17 1 Introduction

18 Offline multi-agent reinforcement learning (MARL) aims to learn coordinated policies from a fixed
19 dataset, without further environment interaction. The core difficulty is that coordination is defined
20 over the *joint action* distribution of all agents: individually plausible per-agent actions may combine
21 into unseen joint actions that violate safety or yield low value. This is especially problematic when
22 the data contains multiple valid coordination modes, defined as distinct joint behavior patterns, where
23 a unimodal policy averages alternatives into an invalid in-between action. Expressive generative
24 policies based on diffusion and flow matching address this multimodality by modeling joint action as
25 a multi-mode distribution [1, 2, 3], but their multi-step sampling can be costly at decision time and
26 complicate value-guided actor extraction through critic gradients [4]. A recent line of work therefore
27 trains a centralized generative *teacher* over joint actions and *distills* it into decentralized one-step
28 *students*, which act locally and admit more direct value-guided extraction [2].

29 Distillation, however, exposes a gap between joint modeling and decentralized execution. The teacher
30 conditions on the full joint context and a shared noise sample, whereas each local actor sees only its
31 own context and independent noise. If the teacher’s choice between coordination modes depends on
32 information hidden from a local actor, the ℓ_2 distillation loss regresses to the conditional mean and
33 collapses distinct modes into an invalid in-between action. Mode-preserving distillation therefore
34 requires not only that the teacher cover the target modes, but also that those modes be *compatible*
35 *with the information available to local actors*. This makes coherent teacher training critical: standard
36 flow-based training pairs noise samples with replay targets independently, leaving an unstructured
37 mapping between noise and joint behaviors. With multimodal targets, adjacent noise points are




Figure 1: **Overview of source-to-mode routing.** (Left) **Standard Flow Policy:** independent noise-target pairing can entangle nearby source samples with incompatible coordination modes, producing off-support fan mass. (Middle) **MoSDOT:** semi-discrete optimal transport partitions noise space into capacity-matched cells assigned to coordinated behavior modes. (Right) Route quality is measured by off-support avoidance (*fan-free endpoints*), consistency of source-to-mode routing (*clean route consistency*), endpoint accuracy to valid modes (*strict success*), and balanced mass across modes (*balanced coverage*). Full metric definitions are provided in the appendix.

forced to learn conflicting modes simultaneously, producing trajectories that pass through low-density regions or form spurious bridges between modes (Fig. 1, left). Because independent local policies capture only the conditional mean, even minor mode-mixing in the teacher can lead the student to collapse onto invalid joint actions. A clean teacher alone, however, is not sufficient. Independent agents cannot in general represent correlated joint support such as the anti-correlated XOR target, a residual mismatch we call the *product-projection gap*.

Our *Mode-Support Semi-Discrete Optimal Transport* (MoSDOT) addresses the teacher-side problem by summarizing multimodal replay into a finite set of representative joint behavior modes — the *mode support* — with prescribed capacities, and partitioning the noise space via semi-discrete optimal transport into cells assigned to those modes before teacher training (Fig. 1, middle). MoSDOT thus reduces teacher-side routing artifacts and preserves the mode support that local actors can recover, while leaving the product-projection gap explicit. To investigate this residual gap, we additionally consider a common-randomness variant in which a portion of the noise input acts as an execution-time shared signal: all agents read the same coordination mode from this signal but still compute actions from their own observations and local noise, allowing correlated joint behavior without exchanging observations or actions.

In controlled diagnostics, MoSDOT mitigates teacher-side routing artifacts and improves source-to-mode assignment consistency. Fig. 1 (right) summarizes these gains across endpoint quality, route consistency, strict success, and coverage balance. On offline MARL benchmarks, MoSDOT performs competitively with existing generative baselines, particularly on datasets containing diverse joint behaviors. Overall, MoSDOT separates teacher-side source-routing artifacts from the residual product-projection gap, and provides a measurable interface between the centralized training and decentralized execution (CTDE) [5].

2 Related Work

2.1 Generative Policies and Distillation in Offline MARL

Offline reinforcement learning (RL) aims to learn policies from fixed datasets by constraining policy improvement close to actions represented in the data to avoid value overestimation [6, 7, 8, 9]. In multi-agent settings, the support is defined over coordinated *joint* actions, so small per-agent deviations can combine into joint behaviors absent from replay. Offline MARL methods address this through implicit constraints [9], actor rectification [10], and global-to-local value regularization under the CTDE. Standard actor or value-factorized parameterizations rarely model multimodal joint action support explicitly, so distinct coordination modes can be averaged into incoherent in-between actions.

Expressive generative policies address this multimodality. Diffusion and flow-based models have been used for trajectory-level planning [11, 12], action generation [13], and one-step offline policy learning [4]. In offline MARL, MADiff [14] models coordinated multi-agent trajectories with diffusion, and DoF [1] factorizes diffusion generation and formulates Individual-Global-identically-Distributed (IGD) generation as a distributional analogue of value factorization. These approaches improve

expressiveness and execution efficiency, but the quality of the centralized generative distribution remains an important prerequisite.

Policy distillation transforms multimodal teachers into efficient one-step students. Since pointwise regression tends to average modes, recent work has introduced set-level objectives based on Implicit Maximum Likelihood Estimation (IMLE) and Chamfer matching [15], as well as drifting objectives that shape support through attraction toward data and repulsion among generated samples [16]. In offline RL, FQL [4] distills a behavioral flow into a Q-regularized one-step policy, and MAC-Flow [2] distills a centralized multi-agent flow teacher into decentralized one-step actors. These methods primarily focus on the distillation process itself, while our work instead studies the upstream construction of the teacher distribution before distillation begins.

2.2 Mode Alignment in Flow Policies and Factorized Execution

Flow-based generative models are often sensitive to how source noise is coupled to data during training. Flow matching trains continuous normalizing flows by regressing vector fields along conditional probability paths [17]. With independent noise–target pairings, multimodal targets can produce samples that fall between valid modes. Optimal transport (OT) flow matching [18] mitigates this effect by coupling noise and data through transport plans, and AlignFlow [19] uses semi-discrete optimal transport (SDOT) to partition source space into Laguerre cells assigned to empirical data samples for image generation.

A related issue arises in MARL, where centralized training must be compatible with decentralized execution. Value factorization methods such as VDN [20], QMIX [21], and QTRAN [22] impose structure on centralized value functions so that decentralized action selection remains aligned with a global objective. Generative offline MARL raises a distributional analogue of this question. A centralized teacher may represent multimodal joint behavior, but the distilled local actors see only local information and are often trained with pointwise regression losses.

Our work focuses on the corresponding upstream failure mode. Before student training, independent teacher pairings can create off-support or entangled source routes between coordinated joint modes. During distillation, mode variation that is not locally observable is further averaged by the student objective. MoSDOT reduces this teacher-side routing artifact by aligning source-noise regions with the mode support via conditional SDOT, providing cleaner teacher targets for support-preserving decentralized distillation.

3 Method

3.1 Problem Setup

We consider cooperative offline MARL under a decentralized partially observable Markov decision process (Dec-POMDP) [23], defined as $\mathcal{M} = \langle \mathcal{N}, \mathcal{S}, \mathcal{O}, \mathcal{A}, P, Z, R, \gamma \rangle$. Let $\mathcal{N} = \{1, \dots, N\}$ be the set of N agents, with $\mathcal{S}$, $\mathcal{O} = \{\mathcal{O}_i\}_{i=1}^N$, and $\mathcal{A} = \{\mathcal{A}_i\}_{i=1}^N$ denoting the joint global state, observation, and action spaces, respectively. We denote $\mathbf{x} = (o_1, \dots, o_N) \in \mathcal{O}$ as the joint observation context and $\mathbf{a} = (a_1, \dots, a_N) \in \mathcal{A}$ as the joint action. $P(s'|s, \mathbf{a}) : \mathcal{S} \times \mathcal{A} \rightarrow \Delta(\mathcal{S})$ is the state transition function and $Z^i(o_i|s) : \mathcal{S} \rightarrow \Delta(\mathcal{O}_i)$ is the i -th agent’s observation function, where $s \in \mathcal{S}$ and $\Delta(\cdot)$ denotes the set of probability measures over the given space. The team reward is generated based on the reward function $R(s, \mathbf{a}, s') : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$, and γ is a discount factor.

At each timestep, all decentralized agents choose their actions $\mathbf{a}$ conditioned on their local observations $\mathbf{x}$ provided by the environment in state s . The environment then transitions to $s' \sim P(s'|s, \mathbf{a})$, and the team reward is given by $R(s, \mathbf{a}, s')$. In the offline setting, learning uses a fixed dataset $\mathcal{D}$ of joint transitions collected by an unknown behavior policy, without further environment interaction.

Our goal is to learn each agent’s policy as a decentralized actor through centralized training. During training, we first use a centralized generative teacher that has access to the joint context $\mathbf{x}$. The teacher is represented by a deterministic sampler conditioned on source noise distribution:

$$\tilde{\mathbf{a}} = \mu_{\psi}(\mathbf{x}, \mathbf{z}), \quad \mathbf{z} = (z_1, \dots, z_N) \sim \rho_0, \quad (1)$$

where ρ_0 is a simple source noise distribution, such as a standard Gaussian. Since $\mathbf{z}$ is random, this sampler induces a conditional distribution over joint actions given $\mathbf{x}$. In the teacher stage, both the




Figure 2: **MoSDOT-based framework for multi-agent distillation.** (a) Replay joint actions are summarized into a mode support $\{y_k(\mathbf{x})\}$ with capacities $\{b_k(\mathbf{x})\}$, and conditional SDOT partitions the noise space into Laguerre cells $\Gamma_{\mathbf{x}}^{-1}(k)$ assigned to each mode. (b) The centralized teacher μ_ψ is distilled into decentralized one-step actors $\{\mu_{w_i}\}$, each using only its local observation and noise. An optional shared signal h correlates local noise to mitigate the strict-product gap.

124 full joint observation context $\mathbf{x}$ and the full joint noise $\mathbf{z}$ may be used. At execution time, however,
 125 each agent must act using only its local information. We therefore deploy a strict-product policy:

$$\pi_{\mathbf{w}}^{\text{prod}}(\mathbf{a} \mid \mathbf{x}) = \prod_{i=1}^N \pi_{w_i}(a_i \mid o_i), \quad (2)$$

126 where $\mathbf{w} = (w_1, \dots, w_N)$ denotes the set of decentralized actor parameters across N agents. In
 127 the one-step reparameterized implementation used throughout this work, the deployed action is
 128 $\hat{\mathbf{a}} = (\mu_{w_1}(o_1, z_1), \dots, \mu_{w_N}(o_N, z_N))$, with each actor receiving only its local observation o_i and
 129 local noise block z_i , unless specified otherwise. Thus, the centralized teacher can model coordinated
 130 joint behavior using $\mathbf{x}$ and $\mathbf{z}$, while the student must reproduce such behavior under this decentralized
 131 information constraint.

132 For value-guided actor extraction, we use a centralized critic $Q_{\text{tot}}(\mathbf{x}, \mathbf{a})$ that satisfies the Individual-
 133 Global-Max (IGM) principle [20, 21, 22]:

$$\arg \max_{\mathbf{a}} Q_{\text{tot}}(\mathbf{x}, \mathbf{a}) = \left(\arg \max_{a_1} Q_1(o_1, a_1), \dots, \arg \max_{a_N} Q_N(o_N, a_N) \right). \quad (3)$$

134 Here, $Q_{\text{tot}}(\mathbf{x}, \mathbf{a})$ denotes the global Q-function over the joint context and joint action, while $Q_i(o_i, a_i)$
 135 denotes the individual utility for agent i . The IGM principle ensures that greedy selection over each
 136 local Q_i remains consistent with the global optimum of Q_{tot} .

137 The remainder of Sec. 3 describes how we train the teacher so that its source noise is assigned to
 138 coordinated joint-action modes, and how we distill the resulting teacher into decentralized actors.

139 3.2 Mode Support and Capacities

140 Learning the teacher sampler $\mu_\psi(\mathbf{x}, \mathbf{z})$ requires pairing source noise $\mathbf{z}$ with target joint actions so that
 141 the induced conditional distribution matches the target behavior at $\mathbf{x}$. A naive approach pairs $\mathbf{z}$ and
 142 the replay target without considering their relationship, mapping them arbitrarily. When the targets
 143 at a given $\mathbf{x}$ contain multiple distinct coordinated joint behaviors, such as two valid joint actions in
 144 a coordination game, random pairing entangles these behaviors in source-noise space. This often
 145 causes nearby noise vectors to map to unrelated joint actions, and the noise mass allocated to each
 146 behavior is determined by random sampling rather than designed intentionally.

147 We instead organize the targets at each $\mathbf{x}$ into a finite *mode support*, a discrete set of representative
 148 joint actions, and assign a *capacity* (a target source-noise mass) to each element of this support before
 149 training. The sampler then learns a structured assignment in which contiguous regions of $\mathbf{z}$ -space
 150 correspond to single modes, with region sizes set by the capacities.

151 **Mode support.** For each conditioning context $\mathbf{x}$, we collect a finite set of $M_{\mathbf{x}}$ target joint actions:

$$\mathcal{T}_{\mathbf{x}} = \{\mathbf{a}^{(m)}(\mathbf{x})\}_{m=1}^{M_{\mathbf{x}}}, \quad \mathbf{a}^{(m)}(\mathbf{x}) \in \mathcal{A}_1 \times \dots \times \mathcal{A}_N. \quad (4)$$

152 In the default construction, these targets are replay actions from datasets. A support aggregation rule

$$q_{\mathbf{x}} : \{1, \dots, M_{\mathbf{x}}\} \rightarrow \{1, \dots, K_{\mathbf{x}}\} \quad (5)$$

groups the targets into $K_{\mathbf{x}}$ clusters. Each cluster k collects the target indices $C_k(\mathbf{x}) = \{m : q_{\mathbf{x}}(m) = k\}$ and is summarized by a representative joint action $\mathbf{y}_k(\mathbf{x}) \in \mathcal{A}_1 \times \dots \times \mathcal{A}_N$, which we call a *mode*. The *mode support* at context $\mathbf{x}$ is the discrete set of these representatives:

$$\mathcal{Y}_{\mathbf{x}} = \{\mathbf{y}_1(\mathbf{x}), \dots, \mathbf{y}_{K_{\mathbf{x}}}(\mathbf{x})\}. \quad (6)$$

The support size $K_{\mathbf{x}}$ is determined by the observed targets and the aggregation rule, not by enumerating the full combinatorial joint action space. The aggregation rule $q_{\mathbf{x}}$ determines the granularity of $\mathcal{Y}_{\mathbf{x}}$. We use identity support for discrete-action benchmarks and a simple joint-action quantization for continuous-action benchmarks.

Capacities. Each mode $\mathbf{y}_k(\mathbf{x})$ is assigned a *capacity* $b_k(\mathbf{x}) \geq 0$ satisfying $\sum_{k=1}^{K_{\mathbf{x}}} b_k(\mathbf{x}) = 1$, specifying its target source-noise mass. We use $b_k(\mathbf{x}) = |C_k(\mathbf{x})|/M_{\mathbf{x}}$ so that the target source-noise mass matches the empirical replay frequency of each mode. Although each mode decomposes as $\mathbf{y}_k(\mathbf{x}) = (y_{k,1}(\mathbf{x}), \dots, y_{k,N}(\mathbf{x}))$, source assignment is performed on $\mathbf{y}_k(\mathbf{x})$ as a whole, so source-noise regions are assigned to coordinated joint behaviors rather than to a product of per-agent action modes. This is what allows the sampler to represent correlated joint behaviors that a per-agent factorization cannot.

3.3 Source Assignment via Semi-Discrete Optimal Transport

Sec. 3.2 defines the target mass $b_k(\mathbf{x})$ for each mode, but leaves open which source-noise samples should be paired with which mode. We resolve this by partitioning source-noise space so that each cell carries the prescribed mass and is assigned to a single mode. We adopt the SDOT-based source–target alignment introduced for image generation by Kong et al. [19], and apply it to the capacity-weighted mode support $\mathcal{Y}_{\mathbf{x}}$.

Let ρ_0 denote the source law over joint noise $\mathbf{z} = (z_1, \dots, z_N)$. We define the discrete target measure at context $\mathbf{x}$ as follows:

$$\nu_{\mathbf{x}} = \sum_{k=1}^{K_{\mathbf{x}}} b_k(\mathbf{x}) \delta_{\mathbf{y}_k(\mathbf{x})}, \quad (7)$$

where $\delta_{\mathbf{y}_k(\mathbf{x})}$ denotes a Dirac measure at mode $\mathbf{y}_k(\mathbf{x})$, so $\nu_{\mathbf{x}}$ places mass $b_k(\mathbf{x})$ on each mode in $\mathcal{Y}_{\mathbf{x}}$. For a cost $c_{\mathbf{x}}(\mathbf{z}, k)$ between source noise $\mathbf{z}$ and mode $\mathbf{y}_k(\mathbf{x})$, the SDOT primal is defined as follows:

$$\min_{\Gamma_{\mathbf{x}}} \mathbb{E}_{\mathbf{z} \sim \rho_0} [c_{\mathbf{x}}(\mathbf{z}, \Gamma_{\mathbf{x}}(\mathbf{z}))] \quad \text{s.t.} \quad \rho_0(\Gamma_{\mathbf{x}}^{-1}(k)) = b_k(\mathbf{x}), \quad k = 1, \dots, K_{\mathbf{x}}. \quad (8)$$

Here, $\Gamma_{\mathbf{x}} : \mathbf{z} \mapsto k$ is the source-to-mode assignment, mapping each source sample to a mode index. We take $c_{\mathbf{x}}$ to be the squared Euclidean distance in the representation that defines $\mathcal{Y}_{\mathbf{x}}$. The SDOT dual is a finite-dimensional concave problem and admits stochastic optimization on samples from ρ_0 [24]. We write its objective as:

$$\Phi_{\mathbf{x}}(g) = \sum_{k=1}^{K_{\mathbf{x}}} b_k(\mathbf{x}) g_k + \mathbb{E}_{\mathbf{z} \sim \rho_0} \left[\min_{j \in \{1, \dots, K_{\mathbf{x}}\}} (c_{\mathbf{x}}(\mathbf{z}, j) - g_j) \right]. \quad (9)$$

The optimal dual weights are obtained as $g_{\mathbf{x}}^* \in \arg \max_{g \in \mathbb{R}^{K_{\mathbf{x}}}} \Phi_{\mathbf{x}}(g)$. The resulting Laguerre cells define the source-to-mode assignment:

$$\Gamma_{\mathbf{x}}(\mathbf{z}) = \arg \min_k [c_{\mathbf{x}}(\mathbf{z}, k) - g_{x,k}^*]. \quad (10)$$

We solve Eq. 9 stochastically with Monte-Carlo sampling from ρ_0 . For each pair $(\mathbf{x}, \mathbf{z})$, the assignment $\Gamma_{\mathbf{x}}$ selects the mode index, and we write the corresponding mode as $\mathbf{y}_{\Gamma}(\mathbf{x}, \mathbf{z}) = \mathbf{y}_{\Gamma_{\mathbf{x}}(\mathbf{z})}(\mathbf{x})$. We then form the linear interpolation path $\bar{\mathbf{a}}_t = (1-t)\mathbf{z} + t\mathbf{y}_{\Gamma}(\mathbf{x}, \mathbf{z})$ with displacement $\mathbf{v}_{\Gamma}(\mathbf{x}, \mathbf{z}) = \mathbf{y}_{\Gamma}(\mathbf{x}, \mathbf{z}) - \mathbf{z}$, where $t \sim \text{Unif}([0, 1])$. Following [17], the teacher learns a velocity field u_{ψ} by regressing onto the assigned displacement:

$$\mathcal{L}_{\text{teacher}}(\psi) = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}, \mathbf{z} \sim \rho_0, t \sim \text{Unif}([0, 1])} \left[\|u_{\psi}(\bar{\mathbf{a}}_t, t, \mathbf{x}) - \mathbf{v}_{\Gamma}(\mathbf{x}, \mathbf{z})\|_2^2 \right]. \quad (11)$$

At inference, starting from $\mathbf{z}$ and integrating the learned field u_{ψ} from $t = 0$ to $t = 1$ defines the teacher sampler $\mu_{\psi}(\mathbf{x}, \mathbf{z})$ used in Sec. 3.1. Unlike independent pairing, this couples source regions to modes with capacity-controlled mass.

3.4 Joint Policy Distillation

The teacher trained in Sec. 3.3 produces coordinated joint-action samples using the full context and joint noise. We distill these rollouts into the strict-product student family in Eq. (2). For a replay context $\mathbf{x} = (o_1, \dots, o_N)$ and source noise $\mathbf{z} = (z_1, \dots, z_N) \sim \rho_0$, the teacher and student actions are defined as:

$$\tilde{\mathbf{a}} = \mu_\psi(\mathbf{x}, \mathbf{z}), \quad \hat{\mathbf{a}} = (\mu_{w_1}(o_1, z_1), \dots, \mu_{w_N}(o_N, z_N)). \quad (12)$$

The distillation loss matches each local actor to the corresponding component of the centralized teacher rollout:

$$\mathcal{L}_{\text{distill}}(\mathbf{w}) = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}, \mathbf{z} \sim \rho_0} \left[\sum_{i=1}^N \|\mu_{w_i}(o_i, z_i) - [\mu_\psi(\mathbf{x}, \mathbf{z})]_i\|_2^2 \right], \quad (13)$$

where $[\cdot]_i$ extracts the action component for agent i . For value-guided actor extraction, we optimize:

$$\mathcal{L}_\pi(\mathbf{w}) = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}, \mathbf{z} \sim \rho_0} \left[-Q_{\text{tot}}(\mathbf{x}, \hat{\mathbf{a}}) + \alpha \sum_{i=1}^N \|\mu_{w_i}(o_i, z_i) - [\mu_\psi(\mathbf{x}, \mathbf{z})]_i\|_2^2 \right], \quad (14)$$

where $\hat{\mathbf{a}}$ is as defined in Eq. (12), Q_{tot} is the IGM-compatible centralized critic of Eq. (3), and α controls the strength of teacher regularization. The assignment $\Gamma_{\mathbf{x}}$ is used only to train the teacher and is not given to the decentralized actors at execution time.

Strict-product gap. MoSDOT improves the centralized teacher, but it does not remove the execution-time factorization constraint. A strict-product actor cannot represent every correlated joint distribution. For example, even when MoSDOT places clean teacher mass on $(+1, -1)$ and $(-1, +1)$, a strict-product student with the same local marginals also assigns mass to unsupported tuples $(+1, +1)$ and $(-1, -1)$: each agent independently samples its own noise, so the two cannot agree on which mode to execute. We call this residual mismatch the *strict-product gap*.

The MoSDOT construction itself suggests a way to probe this gap. Recall that the assignment $\Gamma_{\mathbf{x}}$ partitions the noise space into Laguerre cells, each mapped to a single mode $\mathbf{y}_k(\mathbf{x})$. If all agents observe the same cell index, they automatically agree on the same mode and avoid unsupported product tuples. This motivates a shared-randomness variant. Let $h \sim \rho_h$ be a random signal shared across agents at execution, playing the role of a shared cell index, and let $\epsilon_i \sim \rho_i$ be agent-local noise for within-mode variability. The resulting policy can be modified as:

$$\pi_{\mathbf{w}}^{\text{sr}}(\mathbf{a} \mid \mathbf{x}) = \mathbb{E}_{h \sim \rho_h} \left[\prod_{i=1}^N \pi_{w_i}^{\text{sr}}(a_i \mid o_i, h) \right], \quad (15)$$

with one-step implementation $\hat{\mathbf{a}}^{\text{sr}} = (\mu_{w_1}^{\text{sr}}(o_1, h, \epsilon_1), \dots, \mu_{w_N}^{\text{sr}}(o_N, h, \epsilon_N))$. Without h , this reduces to strict-product execution as in Eq. 2. With h , agents act locally and exchange no observations, actions, or messages, but their mode choice is correlated through the shared signal. This isolates the residual error due to independent local sampling and complements the teacher-side coupling provided by MoSDOT. Further details are provided in Appendix.

4 Experiments

Our experiments separate three sources of error in CTDE generative distillation: (i) teacher-side off-support mass, (ii) incompatibility between the teacher’s source assignment and decentralized actors, and (iii) the residual strict-product gap. MoSDOT is designed to address first two by replacing independent source–target pairing with a capacity-matched SDOT assignment. We then compare strict-product and shared-randomness decentralized actors to isolate what can be preserved under decentralized execution.

4.1 Experimental Setup

Diagnostics and benchmarks. We use two controlled tasks to visualize the failure modes studied in this work. The landmark diagnostic in Fig. 3 evaluates teacher-side source assignment with multiple coordinated modes across three agents. The XOR task in Fig. 6, adapted from MAC-Flow [2], isolates the strict-product gap using the anti-correlated support $\{(-1, +1), (+1, -1)\}$. For benchmark evaluation, we use SMACv1 [25], SMACv2 [26], and MPE [27].




Figure 3: **Landmark policy trajectories.** Black $\times$ marks denote the six anchor mode endpoints, and colors indicate the three anchor pairs. The top row shows centralized joint teacher rollouts, and the bottom row shows one-step factorized students distilled under independent local noise. Baseline students show fan mass between modes, while MoSDOT reduces fan mass and retains the structure.

Method	Fan-free $\uparrow$	Route cons. $\uparrow$	Success $\uparrow$	Balance $\uparrow$
Flow matching	0.857	0.856	0.915	0.917
IMLE	0.991	0.857	0.848	0.761
Drifting	0.952	0.653	0.985	0.936
MoSDOT (ours)	0.983	0.972	0.991	0.961

Table 1: **Joint-teacher landmark diagnostics.** Mean over 4 seeds.




Figure 4: **Landmark route-locality diagnostics.**

Baselines and protocol. For controlled tasks, we compare against independent flow matching [17], IMLE policy [28], and Drifting model [16]. For benchmark evaluation, we follow the protocol and baseline suite of MAC-Flow [2], which includes representative Gaussian, diffusion, and flow-policy baselines such as MADiff, DoF, and MAC-Flow [14, 1, 2]. Unless stated otherwise, benchmark results use 6 seeds and report two standard deviations after $\pm$. More details on baselines and implementations are provided in Appendix E.

Diagnostic metrics. For the landmark diagnostic, we evaluate teacher quality along two axes. **(i) Endpoint support:** does the teacher place samples on the modes in $\mathcal{Y}_x$? *Fan-free* counts samples inside accepted regions around modes. *Success* measures endpoint hits, and *Balance* measures agreement with the target capacities $\{b_k(\mathbf{x})\}$. **(ii) Source routing:** how does the teacher use source noise to choose among modes? *Route consistency* checks whether nearby source samples within an agent’s noise slice receive stable mode labels. *Jacobian off-diagonal ratio* measures how strongly the teacher’s output for agent i depends on other agents’ noise $z_{j \neq i}$ relative to its own noise z_i , with values near zero indicating that mode choice is driven only by locally available information. Endpoint metrics test whether the teacher reaches $\mathcal{Y}_x$, and routing metrics test whether mode choices depend only on locally observable information. Full definitions are given in Appendix E.6.

4.2 Endpoint Accuracy Does Not Imply Distillable Assignments

We first test whether reaching the correct modes is sufficient for factorized distillation. Tab. 1 measures endpoint-level support quality, while Fig. 4 asks a complementary question about whether the teacher’s mode choice for an agent depends on the agent’s own noise or the noise of other agents.

A teacher can reach valid modes while choosing among them in a way unavailable to the corresponding local actor. Flow matching has reasonable route consistency and balance, but its lower fan-free score reveals samples between modes. IMLE achieves the best fan-free score with sharp endpoints, but sacrifices success and balance, appearing as uneven mass across modes in Fig. 4. Drifting shows this mismatch most strongly. It reaches valid modes often, but has the lowest route consistency in Tab. 1 and the largest Jacobian off-diagonal ratio in Fig. 4, indicating its per-agent output is strongly shaped by other agents’ noise—information unavailable to the local actor. Endpoint support is therefore insufficient: when mode choice depends on off-agent noise, distillation can average or mix modes even with accurate endpoints. We evaluate teacher quality along both axes—whether it reaches $\mathcal{Y}_x$, and whether mode choices are reproducible from locally available information.




Figure 6: **XOR strict-product visualization.** (a) Target support placing equal mass on $(-1, +1)$ and $(+1, -1)$. (b) **MAC-Flow**: bridge mass spreads across all quadrants under strict-product execution. (c) **Ours w/o shared signal**: clean teacher, residual strict-product gap. (d) **Ours with shared signal**: gap closed. Blue: factored student; pink: joint teacher.

4.3 MoSDOT Preserves Support Without Entangling Assignments

The previous section showed that no baseline is strong on both axes simultaneously. We now evaluate whether MoSDOT closes this gap. On the same landmark replay, MoSDOT replaces independent source–target pairing with the capacity-matched SDOT assignment of Eq. (11), using replay frequencies as mode capacities.

MoSDOT improves support sharpness without sacrificing assignment locality. On the endpoint axis (Tab. 1), it is competitive with the best endpoint-focused baselines on fan-free support and success, and achieves the best balance. The teacher’s joint mode coverage in Fig. 4 (left, hatched) is also the highest across methods. On the routing axis, its route consistency is the highest among all methods, and its Jacobian off-diagonal ratio in Fig. 4 (right) is comparable to flow matching, well below IMLE and Drifting. Fig. 5 visualizes both axes jointly in Agent 1’s noise plane: Flow leaves a fan between modes, IMLE assigns unbalanced area, and Drift mixes the two modes throughout, while MoSDOT yields balanced, fan-less, disentangled regions.




Figure 5: Source-space partitions in Agent 1’s noise plane.

4.4 Cleaner Source Assignments Survive Decentralized Distillation

Teacher-side improvement is useful only if it persists after decentralized distillation. To examine the teacher-to-student transfer, we distill strict-product students from two teachers under the same student architecture, seed list, aggregation rule, and training budget. The baseline student uses an unrouted flow teacher, while the MoSDOT student uses the capacity-matched teacher from Sec. 4.3.

The student appears to track the teacher’s geometry. Fig. 3 (bottom row) shows that the MoSDOT student largely follows the teacher’s ray structure, while baselines show fan mass between modes. Quantitatively, the MoSDOT-distilled student retains higher fan-free, success, and balance scores than baselines, while baseline students lose substantial endpoint quality through distillation (Appendix F.1). These results suggest that improving the centralized source assignment can improve the support preserved by decentralized actors.

4.5 Shared Randomness Isolates the Strict-Product Gap

Even with a clean teacher, the strict-product student of Eq. (2) cannot in general represent correlated joint support. The XOR target $\{(-1, +1), (+1, -1)\}$ illustrates this: independent local sampling that matches the marginals also produces the unsupported tuples $(+1, +1)$ and $(-1, -1)$. We use the shared-randomness variant of Eq. (15) to separate this product-projection error from teacher-side assignment error. Fig. 6 shows the separation across three stages. The unrouted flow teacher places bridge mass between the two valid tuples, which a strict-product student spreads across all four quadrants (b). MoSDOT concentrates the teacher mass on the two valid tuples, yet independent local sampling still mixes the two modes after distillation (c). Shared randomness, with the same MoSDOT teacher, lets the actors agree on a mode-selection signal while still acting locally, and the student concentrates on the valid tuples (d). The transition from (c) to (d)—where only the execution-time signal changes—isolates the strict-product gap from teacher-side assignment failure.

Table 2: **Performance on discrete-action offline MARL benchmarks (SMACv1 and SMACv2).** Mean over 6 seeds, $\pm 2\sigma$. Best and second-best in sage and powder blue.

Benchmark	Scenario	Dataset	Offline / BC baselines			Diffusion policies			Flow / OT policies		
			BC	MABCQ	MACQL	Diffusion BC	MADiff	DoF	Flow BC	MAC-Flow	MoSDOT
SMACv1	3m	Good	16.0 $\pm$ 1.0	3.7 $\pm$ 1.1	19.1 $\pm$ 0.1	19.5 $\pm$ 0.5	19.3 $\pm$ 0.5	19.8 $\pm$ 0.2	20.0 $\pm$ 0.0	19.8 $\pm$ 0.2	20.0 $\pm$ 0.0
		Medium	8.2 $\pm$ 0.8	4.0 $\pm$ 1.0	13.7 $\pm$ 0.3	13.3 $\pm$ 0.7	16.4 $\pm$ 2.6	18.6 $\pm$ 1.2	14.7 $\pm$ 1.5	18.0 $\pm$ 3.2	18.8 $\pm$ 1.5
		Poor	4.4 $\pm$ 0.1	3.4 $\pm$ 1.0	4.2 $\pm$ 0.1	4.2 $\pm$ 0.2	10.3 $\pm$ 6.1	10.9 $\pm$ 1.1	4.5 $\pm$ 0.1	10.6 $\pm$ 2.2	16.4 $\pm$ 4.9
	8m	Good	16.7 $\pm$ 0.4	4.8 $\pm$ 0.6	18.9 $\pm$ 0.9	19.4 $\pm$ 0.5	18.9 $\pm$ 1.1	19.6 $\pm$ 0.3	19.5 $\pm$ 0.2	19.7 $\pm$ 0.3	20.0 $\pm$ 0.0
		Medium	10.7 $\pm$ 0.5	5.6 $\pm$ 0.6	15.5 $\pm$ 1.5	18.6 $\pm$ 0.6	16.8 $\pm$ 1.6	18.6 $\pm$ 0.8	18.2 $\pm$ 0.8	19.4 $\pm$ 0.6	20.0 $\pm$ 0.0
		Poor	5.3 $\pm$ 0.1	3.6 $\pm$ 0.8	7.5 $\pm$ 1.0	4.8 $\pm$ 0.2	9.8 $\pm$ 0.9	12.0 $\pm$ 1.2	4.9 $\pm$ 0.1	11.5 $\pm$ 0.8	10.8 $\pm$ 0.8
	2s3z	Good	18.2 $\pm$ 0.4	7.7 $\pm$ 0.9	17.4 $\pm$ 0.3	18.0 $\pm$ 1.0	15.9 $\pm$ 1.2	18.5 $\pm$ 0.8	19.5 $\pm$ 0.1	19.5 $\pm$ 0.5	20.1 $\pm$ 0.1
		Medium	12.3 $\pm$ 0.7	7.6 $\pm$ 0.7	15.6 $\pm$ 0.4	13.4 $\pm$ 1.4	15.6 $\pm$ 0.3	18.1 $\pm$ 0.9	15.1 $\pm$ 2.0	17.6 $\pm$ 0.6	18.2 $\pm$ 1.0
		Poor	6.7 $\pm$ 0.3	6.6 $\pm$ 0.2	8.4 $\pm$ 0.8	6.2 $\pm$ 1.2	8.5 $\pm$ 1.3	10.0 $\pm$ 1.1	6.9 $\pm$ 0.8	8.5 $\pm$ 0.6	9.4 $\pm$ 0.8
	5m_vs_6m	Good	15.8 $\pm$ 3.6	2.4 $\pm$ 0.4	16.2 $\pm$ 1.6	16.8 $\pm$ 2.3	16.5 $\pm$ 2.8	17.7 $\pm$ 1.1	14.7 $\pm$ 2.1	18.6 $\pm$ 3.5	18.5 $\pm$ 1.9
		Medium	12.4 $\pm$ 0.9	3.8 $\pm$ 0.5	15.1 $\pm$ 2.9	12.5 $\pm$ 2.1	15.2 $\pm$ 2.6	16.2 $\pm$ 0.9	12.8 $\pm$ 0.8	15.6 $\pm$ 1.3	17.9 $\pm$ 1.2
		Poor	7.5 $\pm$ 0.2	3.3 $\pm$ 0.5	10.5 $\pm$ 3.1	8.0 $\pm$ 1.0	8.9 $\pm$ 1.3	10.8 $\pm$ 0.3	7.7 $\pm$ 0.8	9.8 $\pm$ 2.1	12.3 $\pm$ 0.9
	2c_vs_64zg	Good	17.5 $\pm$ 0.4	10.1 $\pm$ 0.2	12.9 $\pm$ 0.2	17.8 $\pm$ 1.3	14.7 $\pm$ 2.2	16.1 $\pm$ 0.8	18.0 $\pm$ 1.3	19.1 $\pm$ 0.8	20.2 $\pm$ 0.6
		Medium	12.5 $\pm$ 0.3	9.9 $\pm$ 0.2	11.6 $\pm$ 0.1	10.5 $\pm$ 1.1	12.8 $\pm$ 1.2	13.9 $\pm$ 0.9	11.8 $\pm$ 2.6	14.9 $\pm$ 4.1	16.4 $\pm$ 1.0
		Poor	9.7 $\pm$ 0.2	9.0 $\pm$ 0.2	10.2 $\pm$ 0.1	10.2 $\pm$ 2.3	10.8 $\pm$ 1.1	11.5 $\pm$ 1.1	10.0 $\pm$ 0.3	11.4 $\pm$ 0.4	11.4 $\pm$ 1.4
	Average rewards		12.2 $\pm$	5.5 $\pm$	13.1 $\pm$	13.0 $\pm$	13.8 $\pm$	15.6	13.4 $\pm$	15.6	16.7
SMACv2	terrari_5_vs_5	replay	7.3 $\pm$ 1.0	13.8 $\pm$ 4.4	11.8 $\pm$ 0.9	9.3 $\pm$ 0.9	13.3 $\pm$ 1.8	15.4 $\pm$ 1.3	8.3 $\pm$ 1.9	16.6 $\pm$ 4.3	16.6 $\pm$ 4.3
	zerg_5_vs_5	replay	6.8 $\pm$ 0.6	10.3 $\pm$ 1.2	10.3 $\pm$ 3.4	8.1 $\pm$ 1.7	10.2 $\pm$ 1.1	12.0 $\pm$ 1.1	4.6 $\pm$ 0.5	9.8 $\pm$ 1.5	9.7 $\pm$ 2.8
	terrari_10_vs_10	replay	7.4 $\pm$ 0.5	12.7 $\pm$ 2.0	11.8 $\pm$ 2.0	5.5 $\pm$ 1.5	13.8 $\pm$ 1.3	14.6 $\pm$ 1.1	5.8 $\pm$ 1.7	13.0 $\pm$ 4.7	9.4 $\pm$ 2.6
	Average rewards		7.2	12.3	11.3	7.6	12.4	14.0	6.2	13.1	11.9

Table 3: **Performance on continuous-action offline MARL benchmarks (MPE Simple Spread).** *M-R* denotes Medium-Replay.

Scenario	Dataset	Offline MARL					Generative policies		
		MATD3BC	MACQL	ICQ	OMAR	OMIGA	MADiff	MAC-Flow	MoSDOT
Spread	Expert	108.3 $\pm$ 3.3	98.2 $\pm$ 5.2	114.9 $\pm$ 2.6	104.0 $\pm$ 3.4	80.8 $\pm$ 13.8	95.0 $\pm$ 5.3	101.7 $\pm$ 10.9	104.40 $\pm$ 19.66
	Medium	29.3 $\pm$ 4.8	34.1 $\pm$ 7.2	47.9 $\pm$ 18.9	29.3 $\pm$ 5.5	30.1 $\pm$ 16.9	64.9 $\pm$ 7.7	80.1 $\pm$ 20.6	95.97 $\pm$ 12.69
	M-R	15.4 $\pm$ 5.6	20.0 $\pm$ 8.4	37.9 $\pm$ 12.3	13.6 $\pm$ 5.7	5.4 $\pm$ 11.0	30.3 $\pm$ 2.5	50.4 $\pm$ 33.2	54.12 $\pm$ 15.29
	Random	9.8 $\pm$ 4.9	24.0 $\pm$ 9.8	34.4 $\pm$ 5.3	6.3 $\pm$ 3.5	-3.8 $\pm$ 12.3	6.9 $\pm$ 3.1	31.1 $\pm$ 6.8	75.92 $\pm$ 23.81
Average rewards		40.7	44.1	58.8	38.3	28.1	49.2	65.8	82.6

4.6 Offline MARL Benchmark Results

We evaluate MoSDOT on three offline MARL benchmarks. Tab. 2 reports discrete-action results on SMACv1/v2, and Tab. 3 reports continuous-action results on MPE Simple Spread. On SMAC, generative policies—particularly flow- and diffusion-based methods—tend to perform better than Gaussian and behavior cloning (BC) baselines, consistent with the view that a centralized generative teacher is a useful starting point. On MPE, MoSDOT shows clearer gains on *Medium* and *Random*, where replay contains a wider range of joint actions per observation, and remains comparable to existing generative baselines on *Expert*, where replay is concentrated on a single coordinated mode.

These observations are consistent with the view that capacity-matched source assignment is most useful when replay admits a discernible finite mode support. The effect is naturally smaller when the data concentrates on a single coordinated mode (*Expert*), or when the joint-action space grows faster than per-context replay coverage so that the identity-style aggregation rule used here approaches sample-level coupling (e.g., `terrari_10_vs_10`). Designing q_x to recover meaningful mode structure in such regimes is a natural direction for future work.

5 Conclusion

We studied multi-agent generative distillation through how source noise is assigned to coordinated joint behaviors during teacher training. Our analysis shows that endpoint coverage of the teacher is not sufficient for factorized distillation: routing artifacts at the teacher stage propagate to decentralized actors through the distillation loss. MoSDOT addresses this by aligning source noise with a capacity-weighted mode support via conditional SDOT, separating teacher-side routing artifacts from the residual product-projection gap that strict-product execution introduces. A central limitation is the strict-product constraint itself: independent local sampling cannot represent arbitrary correlated joint support, and our shared-randomness variant exposes this gap rather than removing it. Quantifying this residual gap, extending mode-support routing to longer-horizon and value-reweighted settings, and studying when shared execution signals are admissible are natural next steps.

327	Appendix Contents
328	A Proofs and Mathematical Derivations
329	• A.1 Projection Induced by Local Distillation
330	• A.2 Block-Locality of Ordinary Conditional Flow Matching
331	• A.3 Support Error and Reward Ceiling
332	• A.4 Strict-Product Gap: A Minimal XOR Example
333	• A.5 Shared Randomness Closes the Strict-Product Gap
334	B Algorithms
335	• B.1 MoSDOT (Strict-Product Variant)
336	• B.2 MoSDOT-SR (Shared-Randomness Variant)
337	• B.3 Reduction Between the Two Variants
338	C Limitations and Discussion
339	D Implementation Details
340	• D.1 Support Aggregation Rule
341	• D.2 SDOT Solver and Finite-Cache Rebalancing
342	• D.3 Teacher Bootstrap for Critic Learning
343	• D.4 Network Architecture and Flow Schedule
344	E Experimental Details
345	• E.1 Benchmark Environments and Datasets
346	• E.2 Baselines
347	• E.3 Training and Optimization Protocol
348	• E.4 Reporting and Statistical Protocol
349	• E.5 Compute Resources
350	• E.6 Evaluation Metric Definitions
351	F Additional Experiments
352	• F.1 Student-Side Landmark Diagnostics
353	• F.2 Full Route-Locality Diagnostics
354	• F.3 Metric Sensitivity Analysis

A Proofs and Mathematical Derivations

This section collects the formal results that justify the design choices of MoSDOT. Appendix A.1 explains why information that is hidden from a local actor is averaged out by ℓ_2 distillation. Appendix A.2 characterizes when conditional flow matching is block-local at the population optimum, and pinpoints the assumptions that MoSDOT deliberately breaks. Appendix A.3 converts off-support fan mass into a value ceiling that even a perfect student cannot escape. Appendix A.4 shows, on the canonical XOR example, that a clean teacher alone cannot remove the strict-product gap, motivating the shared-randomness variant of Section 3. Appendix A.5 then proves the positive direction: shared randomness produces a mixture of product distributions that exactly represents the XOR target, and characterizes the representational range of this family.

A.1 Projection Induced by Local Distillation

This section formalizes the claim that the ℓ_2 distillation loss regresses each local actor onto the conditional mean of the teacher’s output given local input. Concretely, we show that the population-optimal local predictor satisfies:

$$f_i^*(o_i, z_i) = \mathbb{E}[\mu_\psi(\mathbf{x}, z) \mid o_i, z_i], \quad (16)$$

so that any component of the teacher’s branch decision that remains random after conditioning on (o_i, z_i) is averaged out by the decentralized actor. For any measurable local predictor f , the squared distillation risk decomposes as:

$$\mathbb{E}[\|Y - f(S)\|_2^2] = \mathbb{E}[\|Y - \mathbb{E}[Y \mid S]\|_2^2] + \mathbb{E}[\|\mathbb{E}[Y \mid S] - f(S)\|_2^2]. \quad (17)$$

The first term is independent of f , and the second term is uniquely minimized in L^2 by:

$$f^*(S) = \mathbb{E}[Y \mid S].$$

Substituting $Y = g_i(O, z_{1:N})$ and $S = (o_i, z_i)$ gives:

$$f_i^*(o_i, z_i) = \mathbb{E}[g_i(O, z_{1:N}) \mid o_i, z_i].$$

Therefore, any component of the teacher branch decision that remains random after conditioning on the local input is averaged by the decentralized actor. This is the precise sense in which source assignment affects distillation compatibility: a teacher branch whose mode label depends on z_{-i} cannot be recovered by a strict-product local actor that observes only z_i .

A.2 Block-Locality of Ordinary Conditional Flow Matching

This section clarifies when an ordinary conditional flow-matching teacher is block-local at the population optimum. Let

$$X_i = (1 - T_i)Z_i + T_i A_i, \quad V_i = A_i - Z_i.$$

Assume that, conditioned on the joint context O , the training law factorizes across agents:

$$p(z_{1:N}, a_{1:N}, t_{1:N} \mid O) = \prod_{i=1}^N p(z_i, a_i, t_i \mid O).$$

The squared-loss optimal joint flow field is

$$u_i^*(O, X_{1:N}, T_{1:N}) = \mathbb{E}[V_i \mid O, X_{1:N}, T_{1:N}].$$

Under the conditional factorization above, (V_i, X_i, T_i) is conditionally independent of the off-agent tuples (X_{-i}, T_{-i}) given O , so

$$u_i^*(O, X_{1:N}, T_{1:N}) = \mathbb{E}[V_i \mid O, X_i, T_i].$$

Although a joint network may have access to off-agent chunks, the population MSE optimum derives no benefit from using them under the factorized training law. This condition fails precisely when the teacher objective introduces coupled endpoint assignments, set-level matching, or other gradients that explicitly depend on full joint branches—which is exactly the regime MoSDOT operates in.

389 A.3 Support Error and Reward Ceiling

390 This section formalizes the support-preservation claim used in the main text. For a fixed context x , let
 391 the target mode-support measure be

$$\mu_{\mathbf{x}}^* = \sum_{k=1}^K b_k(\mathbf{x}) \delta_{y_k(\mathbf{x})}.$$

392 An exact source assignment satisfies

$$\Gamma_{x\#} \rho_0 = \sum_{k=1}^K \rho_0(\Gamma_{\mathbf{x}}^{-1}(k)) \delta_{y_k(\mathbf{x})} = \mu_{\mathbf{x}}^*.$$

393 In the ideal assignment limit the source pushforward has zero mass outside the prescribed support
 394 and matches the target capacities exactly.

395 Now let ν be a generated teacher distribution. Let B_k denote the evaluation region associated with
 396 support element $y_k(\mathbf{x})$, and let F denote the off-support fan or bridge region, disjoint from all B_k .
 397 Define

$$m_k = \nu(B_k), \quad m_F = \nu(F).$$

398 The support error on the quotient partition is

$$D_{\text{supp}}(\nu, \mu_{\mathbf{x}}^*) = \frac{1}{2} \left(\sum_{k=1}^K |m_k - b_k(\mathbf{x})| + m_F \right),$$

399 which separates capacity error from off-support mass. If a teacher has small mode-mass error and
 400 small fan mass, it is close to the target support distribution under this quotient evaluation.

401 Assume a bounded reward or value proxy R such that supported modes are optimal and the fan region
 402 has a reward gap $\Delta > 0$:

$$R(y_k(\mathbf{x})) = R^*, \quad R(a) \leq R^* - \Delta \quad \forall a \in F.$$

403 Then any teacher distribution ν with fan mass m_F satisfies

$$J_R(\nu) = \mathbb{E}_{\nu}[R(a)] \leq R^* - \Delta m_F.$$

404 If a student π_S is within total variation distance ϵ of ν , then $\pi_S(F) \geq (m_F - \epsilon)_+$, and therefore

$$J_R(\pi_S) \leq R^* - \Delta(m_F - \epsilon)_+.$$

405 Hence observed off-support mass imposes a value ceiling on faithful distillation whenever that mass
 406 lies in a lower-value region.

407 A.4 Strict-Product Gap: A Minimal XOR Example

408 This section illustrates why a clean teacher does not imply exact preservation by a strict-product
 409 student. Consider two agents and the target anti-correlated support

$$\mathcal{M} = \{(-1, +1), (+1, -1)\}, \quad \mu^* = \frac{1}{2} \delta_{(-1, +1)} + \frac{1}{2} \delta_{(+1, -1)}.$$

410 Let the support reward be $R(x, y) = \mathbf{1}[(x, y) \in \mathcal{M}]$. A strict-product student matching the marginals
 411 of μ^* has

$$q_{\text{clean}}(x, y) = q_x(x) q_y(y), \quad q_x(\pm 1) = q_y(\pm 1) = \frac{1}{2},$$

412 so

$$J_R(q_{\text{clean}}) = q_{\text{clean}}(-1, +1) + q_{\text{clean}}(+1, -1) = \frac{1}{2}.$$

413 Even a clean non-product teacher cannot be exactly represented by independent local sampling.

414 Now consider a fan-contaminated teacher

$$\nu_m = \frac{1-m}{2} \delta_{(-1, +1)} + \frac{1-m}{2} \delta_{(+1, -1)} + m \delta_{(0, 0)}.$$

415 A strict-product student matching the teacher's marginals independently has $q_x(\pm 1) = q_y(\pm 1) =$
 416 $(1-m)/2$ and $q_x(0) = q_y(0) = m$, with support reward

$$J_R(q_m) = \left(\frac{1-m}{2}\right)^2 + \left(\frac{1-m}{2}\right)^2 = \frac{(1-m)^2}{2}.$$

417 Therefore

$$J_R(q_{\text{clean}}) - J_R(q_m) = \frac{1}{2} - \frac{(1-m)^2}{2} = m - \frac{m^2}{2} > 0 \quad \text{for } 0 < m \leq 1.$$

418 Removing fan mass does not eliminate the intrinsic product gap, but it strictly improves the marginally
 419 distilled product student when fan mass contaminates the local marginals. This is the analytic
 420 counterpart of the empirical XOR diagnostic in Fig. 6; the complementary positive result—that the
 421 shared-randomness extension (Algorithm 2) does close this gap—is given in Appendix A.5.

422 A.5 Shared Randomness Closes the Strict-Product Gap

423 Appendix A.4 shows that a strict-product student of the form Eq. (2) cannot represent the anti-
 424 correlated XOR target $\mu^* = \frac{1}{2}\delta_{(-1,+1)} + \frac{1}{2}\delta_{(+1,-1)}$. This section proves the complementary positive
 425 direction for the shared-randomness policy of Eq. (15): a public shared signal h turns the policy into
 426 a *mixture of product distributions*, and this mixture exactly represents anti-correlated supports such
 427 as XOR. We first derive the mixture form from the one-step implementation, then verify exact XOR
 428 reproduction, and finally state the representational range of the resulting policy family.

429 **From one-step implementation to mixture-of-products.** Recall the one-step shared-randomness
 430 sampler $\hat{a}_i = \mu_{w_i}^{\text{sr}}(o_i, h, \epsilon_i)$ with $h \sim \rho_h$ shared across all agents and $\epsilon_i \sim \rho_i$ drawn independently
 431 per agent. We derive Eq. (15) from this implementation in three steps.

432 (i) *Joint conditional law given all noise.* Since the sampler is deterministic given $(h, \epsilon_{1:N})$,

$$p(\mathbf{a} \mid \mathbf{x}, h, \epsilon_{1:N}) = \prod_{i=1}^N \delta(a_i - \mu_{w_i}^{\text{sr}}(o_i, h, \epsilon_i)).$$

433 (ii) *Conditional independence given h .* The local noises $\{\epsilon_i\}$ are mutually independent, so integrating
 434 each ϵ_i against its marginal ρ_i while holding h fixed gives

$$p(\mathbf{a} \mid \mathbf{x}, h) = \prod_{i=1}^N \underbrace{\int \delta(a_i - \mu_{w_i}^{\text{sr}}(o_i, h, \epsilon_i)) \rho_i(d\epsilon_i)}_{=: \pi_{w_i}^{\text{sr}}(a_i \mid o_i, h)} = \prod_{i=1}^N \pi_{w_i}^{\text{sr}}(a_i \mid o_i, h).$$

435 Conditioning on the shared signal restores per-agent independence: the local noises were independent
 436 to begin with, and h is the only cross-agent random variable.

437 (iii) *Marginalization over h .* Integrating the shared signal yields

$$\pi_w^{\text{sr}}(\mathbf{a} \mid \mathbf{x}) = \int p(\mathbf{a} \mid \mathbf{x}, h) \rho_h(dh) = \int \prod_{i=1}^N \pi_{w_i}^{\text{sr}}(a_i \mid o_i, h) \rho_h(dh),$$

438 which is exactly Eq. (15). The integrand is a product over agents, but the integral over h produces a
 439 *mixture* of such products. Conditional independence given h does not imply marginal independence,
 440 so the joint law generally satisfies $a_i \not\perp a_j$ even though $a_i \perp a_j \mid h$.

441 **Exact XOR reproduction.** For the target $\mu^* = \frac{1}{2}\delta_{(-1,+1)} + \frac{1}{2}\delta_{(+1,-1)}$ of Appendix A.4, take a
 442 binary shared signal $h \in \{0, 1\}$ with $\rho_h(0) = \rho_h(1) = \frac{1}{2}$ and degenerate local noises ($d_{\epsilon_i} = 0$).
 443 Define the per-agent maps

$$\mu_{w_1}^{\text{sr}}(\cdot, h, \cdot) = \begin{cases} -1 & h = 0, \\ +1 & h = 1, \end{cases} \quad \mu_{w_2}^{\text{sr}}(\cdot, h, \cdot) = \begin{cases} +1 & h = 0, \\ -1 & h = 1. \end{cases}$$

444 Substituting into Eq. (15) gives

$$\pi_w^{\text{sr}} = \frac{1}{2} \delta_{-1} \otimes \delta_{+1} + \frac{1}{2} \delta_{+1} \otimes \delta_{-1} = \mu^*.$$

445 Each agent acts deterministically from its own input alone, yet the joint law is exactly the anti-
 446 correlated target. The only cross-agent channel is the public draw of h ; no observations, actions, or
 447 messages are exchanged.

448 **Representational range.** The construction generalizes to any target that factorizes through a
 449 discrete latent K ,

$$\mu^*(\mathbf{a} \mid \mathbf{x}) = \sum_{k=1}^K \pi_k(\mathbf{x}) \prod_{i=1}^N \mu_{i,k}^*(a_i \mid o_i, \mathbf{x}),$$

450 by taking ρ_h supported on $\{1, \dots, K\}$ with $\rho_h(k) = \pi_k(\mathbf{x})$ and $\pi_{w_i}^{\text{sr}}(\cdot \mid o_i, k) = \mu_{i,k}^*(\cdot \mid o_i, \mathbf{x})$. A
 451 shared dimension of $\lceil \log_2 K \rceil$ bits suffices for any K -component mixture-of-products target, and
 452 richer ρ_h (e.g. $\rho_h = \rho_0$ on $\mathbb{R}^d$) yields a universal class of correlated decentralized policies under
 453 independent local computation and a single public coordination channel.

454 **Connection to the MoSDOT-SR algorithm.** Algorithm 2 instantiates this mixture with the SDOT
 455 cell index $k = \Gamma_c(h)$ as the discrete latent: mixture weights $\pi_k(\mathbf{x})$ become the capacities $b_{c,k}$, and
 456 product factors $\mu_{i,k}^*$ become anchor coordinates $[y_{c,k}]_i$ plus a local residual $\delta_{w_i}(o_i, [h]_i)$. The number
 457 of mixture components is bounded by the mode-support size K_c , so MoSDOT-SR uses exactly
 458 the representational capacity required by the replay-derived support. MoSDOT defines the mixture
 459 components; MoSDOT-SR exposes a public coordinator that selects among them.

460 B Algorithms

461 This section gives full pseudocode for the two MoSDOT variants used in the main text. Algorithm 1
 462 (Appendix B.1) trains a strict-product student under SDOT-aligned teacher pairing. Algorithm 2
 463 (Appendix B.2) extends it by allowing the decentralized actor to condition on a public shared signal,
 464 and is written as an *additive* extension that reuses the centralized stages of Algorithm 1 verbatim.
 465 Appendix B.3 states the formal reduction between the two variants.

466 B.1 MoSDOT (Strict-Product Variant)

Algorithm 1 MoSDOT: mode-support SDOT alignment and decentralized distillation.

Require: offline replay dataset $\mathcal{D}$ with transitions $(\mathbf{x}, \mathbf{o}, \mathbf{a}, r, x', \mathbf{o}')$; joint source law $\rho_0 = \mathcal{N}(0, I_d)$, $d = \sum_{i=1}^N d_i$, source slices $\mathbf{z} = (z_1, \dots, z_N)$, $z_i \in \mathbb{R}^{d_i}$; context/group map $\chi(\mathbf{x})$, aggregation rule q_c , cost c_c , capacities b_c ; teacher field u_ψ , decentralized actors $\{\mu_{w_i}\}_{i=1}^N$, critic Q_{tot} and target $\bar{Q}_{\text{tot}}$.

Ensure: decentralized actors $\{\mu_{w_i}\}_{i=1}^N$.

CONSTRUCT CONTEXT-CONDITIONED MODE SUPPORT

- 1: **for** each context/support group c induced by $\chi(\mathbf{x})$ **do**
- 2: Collect joint actions $\mathcal{A}_c \leftarrow \{\mathbf{a}^{(m)} \in \mathbb{R}^d : \chi(\mathbf{x}^{(m)}) = c, (\mathbf{x}^{(m)}, \mathbf{a}^{(m)}) \in \mathcal{D}\}$.
- 3: Aggregate into modes $C_{c,k} \leftarrow \{m : q_c(m) = k\}$, $\mathcal{Y}_c \leftarrow \{\mathbf{y}_{c,k}\}_{k=1}^{K_c}$.
- 4: Set normalized capacities $b_{c,k} \geq 0$, $\sum_k b_{c,k} = 1$ *empirical, balanced, or value-weighted.*
- 5: **end for**

COMPUTE SEMI-DISCRETE OT SOURCE ASSIGNMENT

- 6: **for** each context/support group c **do**
- 7: Define target mode law $\nu_c \leftarrow \sum_{k=1}^{K_c} b_{c,k} \delta_{y_{c,k}}$.
- 8: Draw source cloud $\mathcal{Z}_c = \{z^\ell\}_{\ell=1}^L$, $z^\ell \sim \rho_0$, $z^\ell = (z_1^\ell, \dots, z_N^\ell)$.
- 9: Solve the SDOT dual $g_c^* \in \arg \max_{g \in \mathbb{R}^{K_c}} \mathbb{E}_{z \sim \rho_0} [\min_{k \leq K_c} \{c_c(\mathbf{z}, \mathbf{y}_{c,k}) - g_k\}] + \sum_{k=1}^{K_c} b_{c,k} g_k$.
- 10: Define the Laguerre assignment and Monge target lookup $\Gamma_c(z) \leftarrow \arg \min_{k \leq K_c} \{c_c(z, y_{c,k}) - g_{c,k}^*\}$,
 $T_c(z) \leftarrow \mathbf{y}_{c, \Gamma_c(z)}$.
- 11: Optionally rebalance finite cached assignments so that $|\{\ell : \Gamma_c(z^\ell) = k\}|/L \approx b_{c,k}$ (Appendix D.2).
- 12: **end for**
- 13: **while** not converged **do**
- 14: Sample replay batch $\mathcal{B} = \{(\mathbf{x}_b, \mathbf{o}_b, \mathbf{a}_b, r_b, x'_b, \mathbf{o}'_b)\}_{b=1}^B \sim \mathcal{D}$.

TRAIN MOSDOT TEACHER ON SDOT-PAIRED SOURCE-TARGET PAIRS

- 15: Sample alignment contexts $c_b = \chi(\mathbf{x}_b)$ and joint sources $\mathbf{z}_b = (z_{b,1}, \dots, z_{b,N}) \sim \rho_0$.
- 16: Assign each source to its SDOT mode $k_b \leftarrow \Gamma_{c_b}(\mathbf{z}_b)$, $\mathbf{y}_b \leftarrow T_{c_b}(\mathbf{z}_b) = \mathbf{y}_{c_b, k_b}$.
- 17: Sample $t_b \sim \text{Unif}([0, 1])$ and form $\mathbf{a}_{t_b} \leftarrow (1 - t_b)\mathbf{z}_b + t_b\mathbf{y}_b$, $v_b \leftarrow \mathbf{y}_b - \mathbf{z}_b$.
- 18: Update teacher field u_ψ via $\mathcal{L}_{\text{MoSDOT}}(\psi) = \frac{1}{B} \sum_b \|u_\psi(\mathbf{x}_b, \mathbf{a}_{t_b}, t_b) - v_b\|_2^2$.
- 19: Define one-step or integrated teacher action $\mu_\psi(\mathbf{x}, \mathbf{z}) \leftarrow \text{Flow}_{u_\psi}(\mathbf{x}, \mathbf{z})$.

TRAIN CENTRALIZED CRITIC WITH TEACHER BOOTSTRAP

- 20: Sample $\xi'_b \sim \rho_0$, bootstrap action $\mathbf{a}'_b \leftarrow \mu_\psi(\mathbf{x}'_b, \xi'_b)$ *(see Appendix D.3).*
- 21: Update Q_{tot} by $\mathcal{L}_Q = \frac{1}{B} \sum_b (Q_{\text{tot}}(\mathbf{x}_b, \mathbf{a}_b) - r_b - \gamma(1 - d_b)\bar{Q}_{\text{tot}}(\mathbf{x}'_b, \mathbf{a}'_b))^2$.
- 22: Soft-update $\bar{Q}_{\text{tot}} \leftarrow \tau Q_{\text{tot}} + (1 - \tau)\bar{Q}_{\text{tot}}$.

DISTILL DECENTRALIZED ACTORS

- 23: Sample fresh joint source $\zeta_b = (\zeta_{b,1}, \dots, \zeta_{b,N}) \sim \rho_0$.
 - 24: Compute teacher and student actions $\tilde{\mathbf{a}}_b \leftarrow \mu_\psi(\mathbf{x}_b, \zeta_b)$, $\hat{\mathbf{a}}_b \leftarrow (\mu_{w_1}(o_{b,1}, \zeta_{b,1}), \dots, \mu_{w_N}(o_{b,N}, \zeta_{b,N}))$.
 - 25: Update $\{w_i\}_{i=1}^N$ by $\mathcal{L}_{\text{actor}} = \frac{1}{B} \sum_b \|\hat{\mathbf{a}}_b - \tilde{\mathbf{a}}_b\|_2^2 - \lambda_Q \frac{1}{B} \sum_b Q_{\text{tot}}(\mathbf{x}_b, \hat{\mathbf{a}}_b)$.
 - 26: **end while**
 - 27: **return** decentralized actors $\{\mu_{w_i}\}_{i=1}^N$.
-

467 B.2 MoSDOT-SR (Shared-Randomness Variant)

468 Algorithm 2 extends Algorithm 1 by allowing the decentralized actor to condition on a public shared
 469 signal h in addition to local noise ϵ_i , matching the abstract form of Eq. (15). The teacher, SDOT solver,
 470 and critic remain centralized and operate on the joint source $z \sim \rho_0$ exactly as in Algorithm 1; the
 471 shared-signal/local-noise distinction enters only at distillation, where the joint source is reinterpreted
 472 through the factorization $\rho_0 \equiv \rho_h \otimes \rho_1 \otimes \dots \otimes \rho_N$. The pseudocode is stated generically: any actor
 473 parameterization $\mu_{w_i}^{\text{sr}}(o_i, h, \epsilon_i)$ that respects this signature is admissible.

Algorithm 2 MoSDOT-SR: shared-randomness extension of Algorithm 1.

Require: All inputs of Algorithm 1, with the joint source law factorized as $\rho_0 = \rho_h \otimes \bigotimes_{i=1}^N \rho_i$; decentralized
 actors $\{\mu_{w_i}^{\text{sr}}(o_i, h, \epsilon_i)\}_{i=1}^N$; public seed generator G .

Ensure: decentralized actors $\{\mu_{w_i}^{\text{sr}}\}_{i=1}^N$.

*Mode support, SDOT assignment, teacher, and critic stages (Algorithm 1, lines 1–22) are unchanged: they
 treat $z \sim \rho_0$ atomically. The factorization is used only below.*

DISTILL DECENTRALIZED ACTORS

- 1: Sample batch $\mathcal{B} \sim \mathcal{D}$, shared signals $h_b \sim \rho_h$, and local noises $\epsilon_{b,i} \sim \rho_i$.
- 2: Compute teacher and student actions

$$\tilde{\mathbf{a}}_b = \mu_\psi(\mathbf{x}_b, (h_b, \epsilon_{b,1:N})), \quad \hat{\mathbf{a}}_b = (\mu_{w_i}^{\text{sr}}(o_{b,i}, h_b, \epsilon_{b,i}))_{i=1}^N.$$

- 3: Update $\{w_i\}$ by $\mathcal{L}_{\text{actor}} = \frac{1}{B} \sum_b \|\hat{\mathbf{a}}_b - \tilde{\mathbf{a}}_b\|_2^2 - \lambda_Q \frac{1}{B} \sum_b Q_{\text{tot}}(\mathbf{x}_b, \hat{\mathbf{a}}_b)$.

DECENTRALIZED EXECUTION

- 4: **for** execution step n **do**
 - 5: All agents publicly reconstruct $h_n = G(s_n) \sim \rho_h$; each agent i privately draws $\epsilon_{i,n} \sim \rho_i$.
 - 6: Agent i outputs $a_{i,n} = \mu_{w_i}^{\text{sr}}(o_{i,n}, h_n, \epsilon_{i,n})$.
 - 7: **end for**
 - 8: **return** $\{\mu_{w_i}^{\text{sr}}\}_{i=1}^N$.
-

474 **Concrete instantiation (Brenier-polar).** In our experiments we instantiate $\mu_{w_i}^{\text{sr}}$ with the Brenier-
 475 polar form:

$$\mu_{w_i}^{\text{sr}}(o_i, h, \epsilon_i) = [y_{c, \Gamma_c(h)}]_i + \delta_{w_i}(o_i, [h]_i),$$

476 which corresponds to taking $\rho_h = \rho_0$ (the entire teacher source is exposed as the shared signal), ρ_i
 477 degenerate ($d_{\epsilon_i} = 0$, no extra local noise), and reusing the SDOT tables $\{\mathcal{Y}_c, g_c^*\}_c$ that are already
 478 cached during the mode-support stage. The shared seed generator G then publishes a fresh $h_n \sim \rho_h$
 479 at each execution step, every agent reconstructs the same cell index $\Gamma_c(h_n)$ via the public table,
 480 and each agent emits its coordinate $[y_{c, \Gamma_c(h_n)}]_i$ plus a local residual $\delta_{w_i}(o_{i,n}, [h_n]_i)$. Algorithm 2
 481 does not commit to this choice; any actor matching the $\mu_{w_i}^{\text{sr}}(o_i, h, \epsilon_i)$ signature—for example, a fully
 482 MLP actor that concatenates h and ϵ_i , or a partial-coordinator variant with $0 < d_h < \dim \rho_0$ —is
 483 admissible.

484 B.3 Reduction Between the Two Variants

485 Algorithm 1 (strict-product) is recovered from Algorithm 2 (MoSDOT-SR) precisely when the per-
 486 agent actor $\mu_{w_i}^{\text{sr}}(o_i, h, \epsilon_i)$ does not depend on the shared signal h . We give two equivalent operational
 487 characterizations and then position the two algorithms as endpoints of a single design family.

488 **(i) Degenerate shared signal:** $d_h = 0$. Setting the dimension of the shared component to zero
 489 collapses the joint source factorization to $\rho_0 = \rho_1 \otimes \dots \otimes \rho_N$, so every joint draw reduces to
 490 $z = \epsilon_{1:N}$. The SDOT assignment Γ_c becomes a function of $\epsilon_{1:N}$ alone, the public seed generator
 491 G has no work to do (its codomain is a singleton), and at execution every agent samples its own
 492 $\epsilon_i \sim \rho_i$ independently. The student family $\{\mu_{w_i}^{\text{sr}}(o_i, \cdot, \epsilon_i)\}_i$ then depends on local information only,
 493 recovering the strict-product actor of Algorithm 1.

494 **(ii) h -independent actor.** Even when ρ_h is non-degenerate, if $\mu_{w_i}^{\text{sr}}(o_i, h, \epsilon_i) = \mu_{w_i}(o_i, \epsilon_i)$ for all
 495 (h, ϵ_i) , the integral over h in Eq. (15) factorizes:

$$\pi_w^{\text{sr}}(\mathbf{a} \mid \mathbf{x}) = \int \prod_{i=1}^N \pi_{w_i}^{\text{sr}}(a_i \mid o_i, h) \rho_h(dh) = \prod_{i=1}^N \pi_{w_i}(a_i \mid o_i),$$

496 which is exactly the strict-product policy of Eq. (2). The shared signal is still drawn but unused; the
 497 teacher, SDOT, and critic stages of Algorithm 2 run as in Algorithm 1 since they operate on $z \sim \rho_0$
 498 irrespective of how the source decomposes into $(h, \epsilon_{1:N})$.

499 **Endpoints of one family.** Algorithm 2 therefore parameterizes a continuous family of decentralized
 500 policies indexed by how much of the joint teacher source is exposed as a public coordinator:

- 501 • *Strict-product limit* ($d_h = 0$): no coordination signal; recovers Algorithm 1 and inherits the
 502 product-projection gap of Appendix A.4.
- 503 • *Brenier-polar limit* ($\rho_h = \rho_0, d_{\epsilon_i} = 0$): the entire teacher source is shared and the SDOT cell
 504 index acts as a global mode coordinator; the strict-product gap closes on supports representable
 505 by the mode capacities, with the formal expressivity argument given in Appendix A.5.
- 506 • *Intermediate cases* ($0 < d_h < \dim \rho_0$): partial coordinators. The shared component carries
 507 enough randomness to break product factorization while leaving residual local degrees of
 508 freedom; we leave empirical study of this regime to future work.

509 The two algorithms thus do not stand as separate methods but as endpoints of a single design axis: how
 510 much of the joint teacher source is published as a decentralized coordinator. The teacher, SDOT solver,
 511 and critic are shared between the two endpoints; only the distillation stage and the execution-time
 512 actor signature differ.

513 C Limitations and Discussion

514 We close by collecting the main limitations of MoSDOT and noting partial remedies where they exist.

515 **Strict-product execution is fundamentally limited.** Independent local sampling cannot represent
 516 arbitrary correlated joint support: even with a perfect teacher, the strict-product student of Eq. (2) fails
 517 on anti-correlated targets such as XOR (Appendix A.4). MoSDOT improves the centralized teacher but
 518 does not remove this execution-time constraint. The shared-randomness variant (Appendix A.5) closes
 519 the gap on supports representable as mixtures of products, but requires a public coordination channel,
 520 which is not always admissible in deployment. Quantifying when such a channel is available—and
 521 how much shared dimension is actually needed—remains open.

522 **Empirical scope.** We evaluate on three offline MARL benchmarks (SMACv1, SMACv2, MPE
 523 Simple Spread) with at most 10 agents and either small discrete or low-dimensional continuous action
 524 spaces. Scaling to larger agent populations, richer action spaces, or online fine-tuning is left for future
 525 work, as is direct comparison with communication-augmented baselines that use richer cross-agent
 526 channels than the public seed considered here.

527 D Implementation Details

528 This section documents implementation choices that are specific to MoSDOT itself, as opposed to
 529 general experimental settings that apply to all methods (those are deferred to Appendix E).

530 D.1 Support Aggregation Rule

531 The aggregation rule q_x specifies the granularity of the replay-derived target support and is not treated
 532 as an oracle mode label. The identity rule $q_x(m) = m$ is always admissible and recovers the raw
 533 sample support. When repeated, quantized, or high-confidence clustered joint targets are present, q_x
 534 may merge those targets into a smaller support. We use *conservative aggregation*: ambiguous targets
 535 remain separate rather than being forced into a shared support element. The representation used by

536 q_x can be a joint action, an action chunk, a short-horizon endpoint, or a learned target embedding.
 537 In all cases the support index refers to a full joint target rather than independently chosen per-agent
 538 modes. This convention prevents the assignment problem from degenerating into a product of local
 539 mode assignments. For discrete-action benchmarks we use the identity rule; for continuous-action
 540 benchmarks we use a simple joint-action quantization at fixed bin width.

541 D.2 SDOT Solver and Finite-Cache Rebalancing

542 The Laguerre assignment Γ_x defined in Sec. 3 is specified for the population source law ρ_0 . In
 543 implementation we approximate ρ_0 by a finite Monte-Carlo cache: for each context/support group c
 544 we draw $\mathcal{Z}_c = \{z^\ell\}_{\ell=1}^L \stackrel{\text{i.i.d.}}{\sim} \rho_0$ of size $L = 65,536$ and define the empirical cell masses

$$\hat{m}_{c,k}^{(L)} = \frac{1}{L} |\{\ell : \Gamma_c(z^\ell) = k\}|.$$

545 The cost c_c is the squared Euclidean distance in the joint-action representation that defines $\mathcal{Y}_c$.

546 **Dual solver.** The SDOT dual is a K_c -dimensional concave program. We solve it on the cache with
 547 the damped-Newton scheme of Kitagawa et al. [24], terminating when every empirical cell mass
 548 $\hat{m}_{c,k}^{(L)}$ matches its target capacity $b_{c,k}$ within 10^{-3} relative tolerance.

549 **Finite-cache rebalancing.** With finite L , residual deviations between $\hat{m}_{c,k}^{(L)}$ and $b_{c,k}$ may persist
 550 even at the converged dual. We remove them with the rebalancing step of Kong et al. [19], reassigning
 551 boundary cache samples to neighboring cells with the smallest increase in transport cost until exact
 552 mass-matching is achieved. This step acts only on the finite cache and does not modify the population
 553 SDOT problem; in our landmark caches it reassigns 0.6% to 1.6% of cached entries.

554 D.3 Teacher Bootstrap for Critic Learning

555 In our value-guided implementation, the assignment-trained teacher provides bootstrap actions for
 556 critic learning. For a replay transition (o_i, a_i, r_i, o'_i) we sample a joint next action from the teacher,

$$\mathbf{a}' = \mu_\psi(o', z'), \quad z' \sim \rho_0,$$

557 and train an individual critic with

$$\mathcal{L}_Q(\theta_i) = \mathbb{E} \left[\left(Q_{\theta_i}(o_i, a_i) - (r_i + \gamma(1-d)\bar{Q}_{\bar{\theta}_i}(o'_i, a'_i)) \right)^2 \right].$$

558 The decentralized actors are then updated with the value-guided distillation objective in Eq. (14).
 559 This teacher-bootstrap variant is an implementation choice; the core MoSDOT construction is the
 560 source assignment used to train the centralized teacher.

561 D.4 Network Architecture and Flow Schedule

562 All networks are [512, 512, 512, 512] MLPs with layer normalization. The teacher conditions on the
 563 joint observation and joint noise; the decentralized students take only their local observation and
 564 local noise slice. The MoSDOT-SR variant additionally uses a trainable noise map $r_{\eta_i}(o_i, h, \epsilon_i)$ that
 565 combines the public shared signal h with the local noise ϵ_i . The teacher uses a linear-interpolation
 566 flow with 10 Euler steps for both training-time bootstrap actions and inference; ablations on Euler
 567 step count are reported in Appendix F.

568 E Experimental Details

569 E.1 Benchmark Environments and Datasets

570 We evaluate on three offline MARL benchmarks. The choice of maps, quality regimes, and dataset
 571 sources follows the conventions of Lee et al. [2] and Formanek et al. [29] so that comparator scores
 572 remain directly comparable.

SMACv1. SMACv1 [25] simulates decentralized StarCraft II combat between a built-in AI and learned policies. Agents observe partial local information within a sight range, and the discrete action space includes movement, target attack, stop, and no-op. We evaluate on five maps spanning homogeneous and heterogeneous unit compositions: 3m, 8m, 2s3z, 5m_vs_6m, and 2c_vs_64zg. For each map we use the Good, Medium, and Poor quality datasets from the off-the-grid benchmark [29], giving 15 dataset–scenario pairs.

SMACv2. SMACv2 [26] extends SMACv1 with randomized start positions, randomized unit types, and rebalanced sight and attack ranges, increasing the diversity of episodes. We evaluate on three maps under the Replay regime: terran_5_vs_5, zerg_5_vs_5, and terran_10_vs_10. Datasets are again drawn from the off-the-grid benchmark.

MPE Simple Spread. MPE [27] is a continuous-action particle environment for cooperative coordination. In Simple Spread, agents must cover landmarks while avoiding collisions. We use the four quality regimes (Expert, Medium, Medium-Replay, Random) collected by Pan et al. [10]. Following Lee et al. [2], returns are normalized as

$$100 \times \frac{S - S_{\text{random}}}{S_{\text{expert}} - S_{\text{random}}},$$

with $S_{\text{random}} = 159.8$ and $S_{\text{expert}} = 516.8$.

E.2 Baselines

Diagnostic baselines. On the landmark and XOR diagnostics, we compare against three teacher parameterizations. *Independent flow matching* [17] pairs source noise and replay targets independently within each context. *IMLE policy* [28] uses a set-level objective based on implicit maximum likelihood estimation to improve endpoint coverage and reduce mode collapse. *Drifting* [16] trains generated samples through attraction toward the target support and repulsion among generated samples.

Benchmark baselines. For benchmark evaluation we additionally compare against three categories of offline MARL baselines reported in Lee et al. [2], all reproduced under the same dataset splits and seed protocol.

Gaussian-policy baselines. BC is a behavioral cloning policy parameterized as a Gaussian (continuous) or categorical (discrete) head. MABCQ [30] extends BCQ [6] to multi-agent settings with a Q-mixer. MACQL extends CQL [7] with a conservative penalty. MATD3BC extends TD3+BC [8] with $\alpha = 2.5$. ICQ [31] uses an implicit constraint approach for joint policy regularization. OMAR [10] combines CQL-style critic regularization with a CEM-based policy improvement head. OMIGA [5] uses a state-dependent value mixer with advantage-weighted regression [32].

Diffusion-policy baselines. Diffusion BC is a behavioral cloning policy with a 200-step Gaussian DDPM scheduler. MADiff [14] models coordinated multi-agent trajectories with a centralized-then-decentralized diffusion policy. DoF [1] factorizes the diffusion process across agents under the IGM-compatible W-concat factorization with 200 denoising steps.

Flow-policy baselines. Flow BC trains an independent flow policy per agent with the same 10-step linear-interpolation schedule as MAC-Flow but without the joint teacher and IGM critic. MAC-Flow [2] is the immediate predecessor of MoSDOT and shares the centralized-teacher to decentralized one-step actor pipeline. MoSDOT replaces the independent source–target pairing of MAC-Flow’s teacher with the SDOT-aligned coupling.

All baselines use the same [512, 512, 512, 512] MLP architecture as MoSDOT and follow the official implementations cited above. Public repositories used for reimplementations are listed in the supplementary code release.

E.3 Training and Optimization Protocol

Optimization. We use Adam with learning rate 3×10^{-4} for both actor and critic, batch size 64, weight decay 0, optimizer epsilon 10^{-5} , Polyak averaging coefficient $\tau = 0.005$, and discount factor $\gamma = 0.995$. The BC distillation coefficient is $\alpha = 3.0$ by default.

619 **Training schedule.** We train for 10^6 gradient steps on SMACv1 and SMACv2, and 5×10^5 steps on
 620 MPE, matching the budget of Lee et al. [2]. We evaluate every 50k steps with 10 evaluation episodes.

621 E.4 Reporting and Statistical Protocol

622 For benchmark scores we run 6 random seeds and report mean $\pm$ two standard deviations following
 623 Lee et al. [2]. For diagnostic metrics we report mean $\pm$ one standard deviation over 4 seeds. Each
 624 metric is averaged within a seed across agents and contexts before being aggregated across seeds.
 625 Best-value bolding follows a paired seed-level test at significance level 0.05.

626 E.5 Compute Resources

627 All experiments run on a single NVIDIA RTX 4090 GPU with CUDA 12.2 and JAX 0.4.30. Training
 628 one MoSDOT seed takes approximately the same wall-clock time as one MAC-Flow seed. The SDOT
 629 cache is solved once per context group before teacher training and amortized over the 10^6 gradient
 630 steps.

631 E.6 Evaluation Metric Definitions

632 This subsection defines the diagnostic metrics used in Section 4 and Appendix F. Teacher-side metrics
 633 (Appendix E.6.1) test whether the teacher reaches the mode support, and route-locality metrics
 634 (Appendix E.6.2) test whether mode choices depend only on locally observable information.

635 E.6.1 Teacher-Side Support Metrics

636 **Fan-free.** A generated sample is assigned to the nearest support region if it lies within the prescribed
 637 acceptance region of some support element. Samples outside all support regions are counted as fan
 638 mass. *Fan-free* reports the complementary fraction. Higher is better.

639 **Balance.** Let $\hat{m}_k$ be the sampled mass assigned to support element k . The mode-mass error is

$$\frac{1}{K_{\mathbf{x}}} \sum_{k=1}^{K_{\mathbf{x}}} |\hat{m}_k - b_k(\mathbf{x})|.$$

640 *Balance* reports an inverse transformation of this error so that higher is better.

641 **Success.** *Success* is the fraction of generated samples whose nearest endpoint label matches the
 642 SDOT-assigned mode label. This is a stricter endpoint-accuracy criterion than fan-free.

643 **Route consistency.** Route consistency is computed by applying a common post-hoc support
 644 assignment rule to all methods and measuring whether nearby source samples within an agent’s noise
 645 slice map to stable support labels.

646 **Joint mode coverage.** For a threshold $\tau > 0$, joint mode coverage is the fraction of support
 647 elements with sampled mass at least τ , computed as $\frac{1}{K_{\mathbf{x}}} \sum_{k=1}^{K_{\mathbf{x}}} \mathbf{1}[\hat{m}_k \geq \tau]$.

648 E.6.2 Route-Locality Metrics

649 Route locality is reported through two complementary families. The sample-level family operates on
 650 replayed teacher samples. The model-level family operates on the restored teacher MLP and provides
 651 causal evidence for cross-noise dependence.

652 **Own-cell entropy.** For each agent i , z_i is partitioned into a $b \times b$ quantile grid with default $b = 10$.
 653 Within each cell we compute the empirical entropy of the branch label $m_i \in \{0, 1\}$, and the reported
 654 value is the cell-mass-weighted average. Sensitivity at $b \in \{5, 10, 20\}$ is reported in Appendix F.3.

655 **KNN flip rate.** For each sample we find its k nearest neighbors in z_i with default $k = 20$, and
 656 report the fraction of neighbors with a different branch label m_i . Sensitivity at $k \in \{10, 20, 50\}$ is
 657 reported in Appendix F.3.

Off-diagonal AUC excess. We train a logistic classifier to predict m_i from $z_{j \neq i}$ and report its AUC minus the chance baseline of 0.5. A value near zero indicates that other agents’ noise carries no information about agent i ’s branch.

Jacobian off-diagonal ratio. On the restored teacher MLP, we compute the endpoint Jacobian $J_i = [\partial a_i / \partial z_1, \dots, \partial a_i / \partial z_N]$ at $t = 1$ via backward-mode autodiff. The reported ratio is

$$\frac{\sum_{j \neq i} \|\partial a_i / \partial z_j\|_F}{\|\partial a_i / \partial z_i\|_F},$$

averaged over agents and replayed source samples. The ratio uses the unclipped endpoint helper so that action saturation does not mask cross-noise sensitivity. The student ratio is exactly 0 by architecture and serves as a structural reference.

Counterfactual flip rate. We fix z_i , resample z_{-i} from ρ_0 ten times per sample, and pass each through the restored teacher MLP. The flip rate is the fraction of resamples that change the nearest-endpoint label of agent i . This is the causal counterpart of the KNN flip diagnostic. It tests whether the model itself uses z_{-i} to select agent i ’s branch, rather than whether the population of trained samples happens to mix.

F Additional Experiments

This section collects supplementary numerical results referenced from the main text. Appendices F.1–F.3 report extended diagnostic results on the controlled tasks: student-side endpoint quality, full sample- and model-level route-locality (entanglement) diagnostics, and a sensitivity study showing that the route-locality ranking is invariant to bin and neighborhood choices.

F.1 Student-Side Landmark Diagnostics

Tab. 4 reports endpoint-quality metrics for one-step factorized students distilled from each teacher under the same student architecture, seed list, and training budget. Comparing this table with the corresponding teacher-side numbers in Tab. 1 quantifies the teacher-to-student transfer described in Section 4: the MoSDOT student largely retains the teacher’s fan-free, success, and balance scores, whereas baselines lose substantial endpoint quality through distillation. This is consistent with the projection result of Appendix A.1—only teachers whose mode choice is locally identifiable can be reproduced by a strict-product student.

Method	Fan-free $\uparrow$	Success $\uparrow$	Balance $\uparrow$
Flow matching	0.849	0.846	0.949
IMLE	0.784	0.558	0.621
Drifting	0.364	0.685	0.328
MoSDOT (ours)	0.971	0.971	0.963

Table 4: **One-step student landmark diagnostics.** Mean over 4 seeds; same 8192 noises $\times$ 3 agents per seed as the teacher table. Best and second-best in sage and powder blue. Higher is better for all metrics.

F.2 Full Route-Locality Diagnostics

Table 5 reports the full route-locality diagnostics, combining sample-level and MLP-level measurements. The sample-level columns (own-cell entropy, mix rate, KNN flip, off-diagonal AUC excess) are computed on replayed teacher samples and capture whether nearby source points within an agent’s noise slice receive consistent mode labels. The MLP-level columns (endpoint Jacobian ratio, counterfactual flip rate) are computed by re-evaluating the restored teacher MLP under controlled noise perturbations, providing *causal* evidence that the model itself uses z_{-i} to choose agent i ’s branch rather than this dependence being an artifact of training-sample correlations.

Across both families, MoSDOT is closest to the lowest-entanglement reference (independent flow matching) and substantially below IMLE and Drifting, while simultaneously producing the cleanest




Figure 7: **Joint mode-tuple support, teacher vs. student.** Bars show the fraction of 32,768 samples (fan included in the denominator) falling into each of the 8 joint mode tuples on the 3-agent landmark task; the dashed line marks the zero-fan uniform ideal 1/8. Top row: centralized teacher (joint flow ODE). Bottom row: one-step factorized student distilled from each teacher under matched architecture and budget. MoSDOT preserves both validity and balance through distillation (95.0% $\rightarrow$ 91.7% valid), while IMLE loses half its valid mass (97.3% $\rightarrow$ 47.1%) and Drift collapses entirely to the fan region (86.3% $\rightarrow$ 1.0%). This is the per-mode visualization of the aggregate scores in Tab. 4 and the empirical counterpart of the projection result in Appendix A.1.

endpoints (Table 1). This is the empirical counterpart of Appendix A.2. When the teacher objective explicitly couples joint endpoints (IMLE, Drifting), block-locality is broken and cross-noise dependence appears. When the coupling is performed upstream by SDOT, endpoint quality improves *without* reintroducing cross-noise dependence.

Method	Sample-level				MLP-level	
	$H(m_i z_i \text{ cell})$	Mix rate	KNN flip	AUC excess	Jac. ratio	CF flip
		(lower is better)			(lower is better)	
Flow	0.019 $\pm$ 0.024	0.006 $\pm$ 0.010	0.001 $\pm$ 0.000	0.020 $\pm$ 0.008	0.066 $\pm$ 0.005	0.000 $\pm$ 0.000
IMLE	0.319 $\pm$ 0.109	0.098 $\pm$ 0.036	0.135 $\pm$ 0.051	0.105 $\pm$ 0.042	0.292 $\pm$ 0.055	0.129 $\pm$ 0.050
Drift	0.666 $\pm$ 0.252	0.244 $\pm$ 0.117	0.315 $\pm$ 0.134	0.290 $\pm$ 0.123	0.772 $\pm$ 0.148	0.307 $\pm$ 0.128
Ours	0.064 $\pm$ 0.024	0.020 $\pm$ 0.010	0.011 $\pm$ 0.003	0.012 $\pm$ 0.003	0.046 $\pm$ 0.015	0.002 $\pm$ 0.002

Table 5: Full route-locality diagnostics. Sample-level metrics use 12 points per method (4 seeds $\times$ 3 agents). H denotes own-noise cell entropy with 10 bins, Mix denotes own-cell mix rate, KNN flip uses $k=20$, and AUC excess reports the maximum off-diagonal LDA AUC minus 0.5. MLP-level metrics report 4 points for the Jacobian ratio (4 seeds, agent-averaged) and 12 points for CF flip (4 seeds $\times$ 3 agents). All values are mean $\pm$ std.

F3 Metric Sensitivity Analysis

Table 6 reports own-cell entropy and KNN flip across bin counts $b \in \{5, 10, 20\}$ and neighborhood sizes $k \in \{10, 20, 50\}$ for the four diagnostic comparators. The relative ordering Flow $\leq$ Ours $\ll$ IMLE $<$ Drift is preserved across all hyperparameter settings, ruling out the hypothesis that the route-locality result depends on a single grid choice.

Method	$H, b=5$	$H, b=10$	$H, b=20$	KNN $k=10$	KNN $k=20$	KNN $k=50$
Flow	0.109 ± 0.061	0.021 ± 0.024	0.005 ± 0.006	0.000 ± 0.000	0.001 ± 0.000	0.001 ± 0.001
IMLE	0.360 ± 0.101	0.319 ± 0.109	0.291 ± 0.106	0.134 ± 0.051	0.135 ± 0.051	0.137 ± 0.051
Drift	0.690 ± 0.240	0.666 ± 0.252	0.634 ± 0.252	0.314 ± 0.134	0.315 ± 0.134	0.316 ± 0.133
Ours	0.181 ± 0.015	0.064 ± 0.024	0.032 ± 0.010	0.008 ± 0.003	0.011 ± 0.003	0.019 ± 0.003

Table 6: Route-locality sensitivity. Each entry reports 12 points per method (4 seeds $\times$ 3 agents), mean $\pm$ std. The relative ranking is preserved across all bin counts and neighborhood sizes.

References

- [1] Chao Li, Ziwei Deng, Chenxing Lin, Wenqi Chen, Yongquan Fu, Weiwan Liu, Chenglu Wen, Cheng Wang, and Siqi Shen. DoF: A diffusion factorization framework for offline multi-agent reinforcement learning. In *ICLR*, pages 34830–34868, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/file/5623c35f3ab5e2c72aeb3abce27dc28f-Paper-Conference.pdf.
- [2] Dongsu Lee, Daehee Lee, and Amy Zhang. Multi-agent coordination via flow matching, 2026. URL <https://arxiv.org/abs/2511.05005>.
- [3] Teng Pang, Zhiqiang Dong, Yan Zhang, Rongjian Xu, Guoqiang Wu, and Yilong Yin. Value-guidance meanflow for offline multi-agent reinforcement learning, 2026. URL <https://arxiv.org/abs/2604.08174>.
- [4] Seohong Park, Qiyang Li, and Sergey Levine. Flow q-learning, 2025. URL <https://arxiv.org/abs/2502.02538>.
- [5] Xiangsen Wang, Haoran Xu, Yinan Zheng, and Xianyu Zhan. Offline multi-agent reinforcement learning with implicit global-to-local value regularization. *NeurIPS*, 36:52413–52429, 2023.
- [6] Scott Fujimoto, David Meger, and Doina Precup. Off-policy deep reinforcement learning without exploration. In *ICML*, pages 2052–2062, 2019.
- [7] Aviral Kumar, Aurick Zhou, George Tucker, and Sergey Levine. Conservative q-learning for offline reinforcement learning. *NeurIPS*, 33:1179–1191, 2020.
- [8] Scott Fujimoto and Shixiang Shane Gu. A minimalist approach to offline reinforcement learning. *NeurIPS*, 34:20132–20145, 2021.
- [9] Ilya Kostrikov, Ashvin Nair, and Sergey Levine. Offline reinforcement learning with implicit q-learning. *arXiv preprint arXiv:2110.06169*, 2021.
- [10] Ling Pan, Longbo Huang, Tengyu Ma, and Huazhe Xu. Plan better amid conservatism: Offline multi-agent reinforcement learning with actor rectification. In *ICML*, pages 17221–17237, 2022.
- [11] Michael Janner, Yilun Du, Joshua B Tenenbaum, and Sergey Levine. Planning with diffusion for flexible behavior synthesis. *arXiv preprint arXiv:2205.09991*, 2022.
- [12] Anurag Ajay, Yilun Du, Abhi Gupta, Joshua Tenenbaum, Tommi Jaakkola, and Pulkit Agrawal. Is conditional generative modeling all you need for decision-making? *arXiv preprint arXiv:2211.15657*, 2022.
- [13] Zhendong Wang, Jonathan J Hunt, and Mingyuan Zhou. Diffusion policies as an expressive policy class for offline reinforcement learning. *arXiv preprint arXiv:2208.06193*, 2022.
- [14] Zhengbang Zhu, Minghuan Liu, Liyuan Mao, Bingyi Kang, Minkai Xu, Yong Yu, Stefano Ermon, and Weinan Zhang. Madiff: Offline multi-agent learning with diffusion models. *NeurIPS*, 37:4177–4206, 2024.

- [15] Ju Dong, Liding Zhang, Lei Zhang, Yu Fu, Kaixin Bai, Zoltan-Csaba Marton, Zhenshan Bing, Zhaopeng Chen, Alois Christian Knoll, and Jianwei Zhang. From flow to one step: Real-time multi-modal trajectory policies via implicit maximum likelihood estimation-based distribution distillation, 2026. URL <https://arxiv.org/abs/2603.09415>.
- [16] Mingyang Deng, He Li, Tianhong Li, Yilun Du, and Kaiming He. Generative modeling via drifting, 2026. URL <https://arxiv.org/abs/2602.04770>.
- [17] Yaron Lipman, Ricky TQ Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching for generative modeling. *arXiv preprint arXiv:2210.02747*, 2022.
- [18] Alexander Tong, Kilian Fatras, Nikolay Malkin, Guillaume Hugué, Yanlei Zhang, Jarrid Rector-Brooks, Guy Wolf, and Yoshua Bengio. Improving and generalizing flow-based generative models with minibatch optimal transport. *arXiv preprint arXiv:2302.00482*, 2023.
- [19] Lingkai Kong, Molei Tao, Yang Liu, Bryan Wang, Jinmiao Fu, Chien-Chih Wang, and Huidong Liu. Alignflow: Improving flow-based generative models with semi-discrete optimal transport, 2025.
- [20] Peter Sunehag, Guy Lever, Audrunas Gruslys, Wojciech Marian Czarnecki, Vinicius Zambaldi, Max Jaderberg, Marc Lanctot, Nicolas Sonnerat, Joel Z Leibo, Karl Tuyls, et al. Value-decomposition networks for cooperative multi-agent learning. *arXiv preprint arXiv:1706.05296*, 2017.
- [21] Tabish Rashid, Mikayel Samvelyan, Christian Schroeder De Witt, Gregory Farquhar, Jakob Foerster, and Shimon Whiteson. Monotonic value function factorisation for deep multi-agent reinforcement learning. *Journal of Machine Learning Research*, 21(178):1–51, 2020.
- [22] Kyunghwan Son, Daewoo Kim, Wan Ju Kang, David Earl Hostallero, and Yung Yi. Qtran: Learning to factorize with transformation for cooperative multi-agent reinforcement learning. In *ICML*, pages 5887–5896, 2019.
- [23] Daniel S Bernstein, Robert Givan, Neil Immerman, and Shlomo Zilberstein. The complexity of decentralized control of markov decision processes. *Mathematics of operations research*, 27(4): 819–840, 2002.
- [24] Jun Kitagawa, Quentin Mérigot, and Boris Thibert. Convergence of a newton algorithm for semi-discrete optimal transport. *Journal of the European Mathematical Society*, 21(9):2603–2651, 2019.
- [25] Mikayel Samvelyan, Tabish Rashid, Christian Schroeder De Witt, Gregory Farquhar, Nantas Nardelli, Tim GJ Rudner, Chia-Man Hung, Philip HS Torr, Jakob Foerster, and Shimon Whiteson. The starcraft multi-agent challenge. *arXiv preprint arXiv:1902.04043*, 2019.
- [26] Benjamin Ellis, Jonathan Cook, Skander Moalla, Mikayel Samvelyan, Mingfei Sun, Anuj Mahajan, Jakob Foerster, and Shimon Whiteson. SMACv2: An improved benchmark for cooperative multi-agent reinforcement learning. *Advances in Neural Information Processing Systems*, 36:37567–37593, 2023.
- [27] Ryan Lowe, Yi I Wu, Aviv Tamar, Jean Harb, OpenAI Pieter Abbeel, and Igor Mordatch. Multi-agent actor-critic for mixed cooperative-competitive environments. *NeurIPS*, 30, 2017.
- [28] Krishan Rana, Robert Lee, David Pershouse, and Niko Sunderhauf. Imle policy: Fast and sample efficient visuomotor policy learning via implicit maximum likelihood estimation, 2025. URL <https://arxiv.org/abs/2502.12371>.
- [29] Claude Formanek, Asad Jeewa, Jonathan Shock, and Arnu Pretorius. Off-the-grid marl: Datasets with baselines for offline multi-agent reinforcement learning. *arXiv preprint arXiv:2302.00521*, 2023.
- [30] Jiechuan Jiang and Zongqing Lu. Offline decentralized multi-agent reinforcement learning. *arXiv preprint arXiv:2108.01832*, 2021.

- 786 [31] Yiqin Yang, Xiaoteng Ma, Chenghao Li, Zewu Zheng, Qiyuan Zhang, Gao Huang, Jun Yang,
787 and Qianchuan Zhao. Believe what you see: Implicit constraint approach for offline multi-agent
788 reinforcement learning. *NeurIPS*, 34:10299–10312, 2021.
- 789 [32] Xue Bin Peng, Aviral Kumar, Grace Zhang, and Sergey Levine. Advantage-weighted regression:
790 Simple and scalable off-policy reinforcement learning. *arXiv preprint arXiv:1910.00177*, 2019.

791 **NeurIPS Paper Checklist**

792 **1. Claims**

793 Question: Do the main claims made in the abstract and introduction accurately reflect the
794 paper’s contributions and scope?

795 Answer: [\[Yes\]](#)

796 Justification: The abstract and introduction state the contributions (MoSDOT teacher-side
797 source assignment, the strict-product gap, and the shared-randomness variant), which are
798 supported by the diagnostics in Sec. 4 and the benchmarks in Sec. 4.6.

799 **2. Limitations**

800 Question: Does the paper discuss the limitations of the work performed by the authors?

801 Answer: [\[Yes\]](#)

802 Justification: Sec. 4 discusses regimes where the effect of MoSDOT diminishes, and we
803 explicitly states the strict-product constraint and remaining open directions as limitations in
804 Appendix.

805 **3. Theory assumptions and proofs**

806 Question: For each theoretical result, does the paper provide the full set of assumptions and
807 a complete (and correct) proof?

808 Answer: [\[Yes\]](#)

809 Justification: All formal results (projection induced by local distillation, block-locality of
810 conditional flow matching, support error and reward ceiling, and the XOR strict-product
811 gap) are stated and proved in Appendix A.

812 **4. Experimental result reproducibility**

813 Question: Does the paper fully disclose all the information needed to reproduce the main ex-
814 perimental results of the paper to the extent that it affects the main claims and/or conclusions
815 of the paper (regardless of whether the code and data are provided or not)?

816 Answer: [\[Yes\]](#)

817 Justification: Algorithms 1 and 2 give full pseudocode, and Appendices C and D specify the
818 architecture, flow schedule, SDOT solver, optimizer, training budget, datasets, and evaluation
819 protocol. The code will also be made publicly available.

820 **5. Open access to data and code**

821 Question: Does the paper provide open access to the data and code, with sufficient instruc-
822 tions to faithfully reproduce the main experimental results, as described in supplemental
823 material?

824 Answer: [\[No\]](#)

825 Justification: Code is not released at submission time. We use publicly available benchmarks
826 (SMACv1, SMACv2, MPE) with the off-the-grid datasets cited in Appendix E.

827 **6. Experimental setting/details**

828 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
829 rameters, how they were chosen, type of optimizer) necessary to understand the results?

830 Answer: [\[Yes\]](#)

831 Justification: Appendix E specifies the optimizer, learning rate, batch size, training schedule,
832 and evaluation protocol.

833	7. Experiment statistical significance
834	Question: Does the paper report error bars suitably and correctly defined or other appropriate
835	information about the statistical significance of the experiments?
836	Answer: [Yes]
837	Justification: Benchmark scores report mean $\pm 2\sigma$ over 6 seeds, and diagnostic metrics report
838	mean $\pm 1\sigma$ over 4 seeds, as described in Appendix E.
839	8. Experiments compute resources
840	Question: For each experiment, does the paper provide sufficient information on the computer
841	resources (type of compute workers, memory, time of execution) needed to reproduce the
842	experiments?
843	Answer: [Yes]
844	Justification: Mentioned in Appendix E.
845	9. Code of ethics
846	Question: Does the research conducted in the paper conform, in every respect, with the
847	NeurIPS Code of Ethics https://neurips.cc/public/EthicsGuidelines ?
848	Answer: [Yes]
849	Justification: The work uses standard simulation benchmarks with no human subjects,
850	personal data, or sensitive content.
851	10. Broader impacts
852	Question: Does the paper discuss both potential positive societal impacts and negative
853	societal impacts of the work performed?
854	Answer: [N/A]
855	Justification: This is foundational research on offline multi-agent RL methodology evaluated
856	on simulation benchmarks, with no direct path to specific societal applications.
857	11. Safeguards
858	Question: Does the paper describe safeguards that have been put in place for responsible
859	release of data or models that have a high risk for misuse (e.g., pre-trained language models,
860	image generators, or scraped datasets)?
861	Answer: [N/A]
862	Justification: The paper does not release pre-trained models, generative content, or scraped
863	data that pose misuse risks.
864	12. Licenses for existing assets
865	Question: Are the creators or original owners of assets (e.g., code, data, models), used in
866	the paper, properly credited and are the license and terms of use explicitly mentioned and
867	properly respected?
868	Answer: [Yes]
869	Justification: All benchmarks (SMACv1, SMACv2, MPE) and baselines are cited in Sec. 4
870	and Appendix E, and used under their original licenses.
871	13. New assets
872	Question: Are new assets introduced in the paper well documented and is the documentation
873	provided alongside the assets?
874	Answer: [N/A]
875	Justification: The paper does not release new datasets or pre-trained models at submission
876	time.
877	14. Crowdsourcing and research with human subjects
878	Question: For crowdsourcing experiments and research with human subjects, does the paper
879	include the full text of instructions given to participants and screenshots, if applicable, as
880	well as details about compensation (if any)?

881 Answer: [N/A]
882 Justification: The paper does not involve crowdsourcing or human subjects.
883 15. **Institutional review board (IRB) approvals or equivalent for research with human**
884 **subjects**
885 Question: Does the paper describe potential risks incurred by study participants, whether
886 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
887 approvals (or an equivalent approval/review based on the requirements of your country or
888 institution) were obtained?
889 Answer: [N/A]
890 Justification: The paper does not involve human subjects.
891 16. **Declaration of LLM usage**
892 Question: Does the paper describe the usage of LLMs if it is an important, original, or
893 non-standard component of the core methods in this research?
894 Answer: [N/A]
895 Justification: LLMs are not part of the core methodology; the proposed method, experiments,
896 and analysis do not involve LLMs.